

# Measuring the Gap: Semantic Alignment Between AI-for-Sustainability Research and SDG Policy Frameworks

## Contents

|          |                                                                        |           |
|----------|------------------------------------------------------------------------|-----------|
| <b>1</b> | <b>Introduction</b>                                                    | <b>3</b>  |
| <b>2</b> | <b>Literature Review</b>                                               | <b>5</b>  |
| 2.1      | AI and the SDGs: What Bibliometrics Shows . . . . .                    | 5         |
| 2.2      | A Three-Level Framework of Research-Policy Alignment . . . . .         | 6         |
| 2.3      | Research-Policy Alignment in AI and Sustainability . . . . .           | 7         |
| 2.4      | Semantic Methods: Rationale and Cautions . . . . .                     | 10        |
| 2.5      | Research Gap . . . . .                                                 | 11        |
| <b>3</b> | <b>Methodology</b>                                                     | <b>11</b> |
| 3.1      | Research Corpus . . . . .                                              | 11        |
| 3.2      | Policy Corpus . . . . .                                                | 13        |
| 3.3      | Embedding Model and Normalisation . . . . .                            | 14        |
| 3.4      | Supervised Reference Classifier . . . . .                              | 14        |
| 3.5      | Scoring and Assignment . . . . .                                       | 16        |
| 3.6      | Classifier Calibration Across Discourse Types . . . . .                | 16        |
| 3.7      | Coverage Gap Analysis . . . . .                                        | 16        |
| 3.8      | Semantic Gap Analysis . . . . .                                        | 18        |
| 3.9      | Coverage–Semantic Interaction . . . . .                                | 19        |
| 3.10     | Key Assumptions and Mitigations . . . . .                              | 19        |
| 3.11     | Methodology Summary . . . . .                                          | 20        |
| <b>4</b> | <b>Results</b>                                                         | <b>22</b> |
| 4.1      | Supervised Reference Classifier Validation . . . . .                   | 22        |
| 4.2      | Coverage Gap: What SDGs Each Corpus Prioritises . . . . .              | 23        |
| 4.3      | Semantic Gap: How Differently Research and Policy Discuss Each SDG . . | 25        |

|          |                                                                                         |           |
|----------|-----------------------------------------------------------------------------------------|-----------|
| 4.4      | Coverage–Semantic Interaction: Two Distinct Dimensions of Observed Divergence . . . . . | 26        |
| 4.5      | Cross-Sensitivity Robustness of the Gap Rankings . . . . .                              | 28        |
| 4.6      | Robustness of the Semantic-Gap Ranking to the Distance Functional . . .                 | 28        |
| <b>5</b> | <b>Discussion</b>                                                                       | <b>32</b> |
| 5.1      | Two Distinct Dimensions of Observed Divergence . . . . .                                | 33        |
| 5.2      | Robust Patterns of Divergence . . . . .                                                 | 34        |
| 5.3      | Implications . . . . .                                                                  | 35        |
| 5.4      | Limitations . . . . .                                                                   | 36        |
| <b>6</b> | <b>Conclusion</b>                                                                       | <b>38</b> |
| <b>A</b> | <b>Diagnostics for Reference Classifier</b>                                             | <b>44</b> |
| A.1      | SDG 4 Lexical Artefact Audit . . . . .                                                  | 44        |
| A.2      | Centroid Similarity . . . . .                                                           | 44        |
| <b>B</b> | <b>Supplementary Robustness and Sensitivity Checks</b>                                  | <b>46</b> |
| B.1      | Lexical Illustration of the Semantic Gap . . . . .                                      | 46        |
| B.2      | Implications for Interpreting the Main Estimates . . . . .                              | 46        |
| <b>C</b> | <b>Sample-Stability Robustness Check</b>                                                | <b>47</b> |
| <b>D</b> | <b>Model Selection: Grid Search Protocol and Rationale</b>                              | <b>48</b> |
| D.1      | Models Tested . . . . .                                                                 | 48        |
| D.2      | LR Grid Search . . . . .                                                                | 49        |
| D.3      | MLP Grid Search . . . . .                                                               | 49        |
| D.4      | Unified Ranking and Final Selection . . . . .                                           | 49        |
| D.5      | Held-Out Test Evaluation of the MLP Robustness Check . . . . .                          | 50        |
| <b>E</b> | <b>Embedding-Model Sensitivity</b>                                                      | <b>50</b> |
| E.1      | Scope . . . . .                                                                         | 50        |
| <b>F</b> | <b>Register-Adjustment Sensitivity</b>                                                  | <b>51</b> |
| <b>G</b> | <b>Concept-Retrieval Sensitivity</b>                                                    | <b>53</b> |
| <b>H</b> | <b>Supplementary Cross-Method Data</b>                                                  | <b>54</b> |
| H.1      | Cross-Method Coverage and Semantic Gap Values . . . . .                                 | 54        |
| H.2      | Detailed Policy Profiles by Source Family . . . . .                                     | 54        |
| <b>I</b> | <b>Declaration of AI Use</b>                                                            | <b>56</b> |

## Abstract

The Sustainable Development Goals (SDGs) are the dominant international development framework. Shared SDG labels do not necessarily imply shared semantic framing between research and policy. Existing bibliometric studies map AI research to SDGs using keyword or classification methods (Armitage et al., 2020; Kashnitsky et al., 2024), but rarely compare the semantic content of research and policy discourse. This study uses Sentence-BERT embeddings and a supervised logistic regression classifier (test macro-F1 = 0.823) to map semantic distance between 3,105,144 AI-for-sustainability abstracts and 40,547 segments from 2,455 institutional SDG policy sources. The analysis separates coverage gaps (which SDGs each corpus emphasises) from within-SDG semantic gaps (how far texts assigned to the same SDG sit apart). Because the corpora differ in register (Biber and Conrad, 2009), these gaps are observed textual-semantic divergence, not substantive disagreement.

The central finding is a dissociation: the coverage gap and the within-SDG semantic gap are independent dimensions of divergence, not linked effects. The primary hypothesis (H1a) tests whether the coverage gap predicts semantic divergence: the association is near-zero in the canonical specification ( $r = -0.083$ ,  $p = 0.751$ ). Research coverage alone likewise does not predict semantic divergence (H1c:  $r = 0.085$ ,  $p = 0.746$ ;  $n = 17$ ). No hypothesis survives robustness checks across alternative assignment methods, policy source families, or segment caps. Together, these null results show that coverage and framing should be measured as separate axes of divergence. The study reports three supporting patterns. First, policy discourse concentrates on SDGs 16 (Peace and Strong Institutions, 35.0%), 13 (Climate Action, 18.8%), and 17 (Partnerships, 23.3%), together 77.1% of policy coverage. Second, research concentrates on SDGs 3 (Good Health, 25.3%), 9 (Industry and Infrastructure, 18.3%), and 4 (Quality Education, artefact-affected, 14.6%). Third, the largest within-SDG semantic gaps occur in SDGs 13 (0.478), 3 (0.464), 11 (0.460), 1 (0.435), and 7 (0.434).

These findings do not prove substantive research-policy misalignment. They show that shared SDG labels are insufficient evidence of shared semantic framing, and that embedding-based distance is a transparent diagnostic map of proximity and divergence. The results need caution because of register effects, policy source-family heterogeneity, OpenAlex retrieval boundaries, and an SDG 4 assignment that is artefact-affected rather than a clean education finding.

## 1 Introduction

Research-policy alignment is often inferred from whether scientific and policy texts address the same topics. Yet shared topical coverage does not establish that those topics are framed in comparable ways. Papers and policy documents can both invoke SDG 13

(Climate Action) or SDG 3 (Good Health) while emphasising different problems, actors, and interventions. High-level thematic agreement can therefore mask deeper divergence in how each discourse community approaches the same goal.

Existing methods for mapping research to the Sustainable Development Goals (SDGs) operate primarily at the level of lexical or bibliometric correspondence — keyword dictionaries, Boolean queries, or citation and mention counts (Armitage et al., 2020; Kashnitsky et al., 2024). These methods measure whether a text is *about* a given SDG, but they cannot distinguish between a text that shares the policy framing of that SDG and one that uses shared vocabulary while addressing entirely different concerns. Coverage and semantic framing are treated as interchangeable, when they may not be.

This paper proposes that research-policy alignment contains at least two separable dimensions — coverage divergence (which SDGs each corpus prioritises) and semantic framing divergence (how differently they discuss the same SDG) — and that these should be measured with independent procedures rather than conflated under shared labels. The paper makes three contributions. First, it advances AI-SDG mapping from Level 1 (lexical correspondence) to Level 2 (observed textual-semantic proximity in embedding space; Section 2.2), treating semantic proximity as a register-sensitive proxy for framing rather than evidence of substantive alignment. Second, it operationalises coverage gaps and semantic gaps separately and measures both across all 17 SDGs using frozen Sentence-BERT embeddings (Reimers & Gurevych, 2019) and a supervised logistic regression classifier trained on pooled, independently annotated reference corpora. Third, it tests whether these two dimensions are related across the 17 SDGs, showing that coverage alone is not a proxy for semantic proximity.

The central research question is: *To what extent do academic AI-for-sustainability research and international institutional SDG policy discourse differ in SDG coverage and semantic framing across the 17 SDGs, and how are these two dimensions related?*

The remainder of this paper is structured as follows. Section 2 reviews the relevant literature, introduces the three-level framework, and positions the study at Level 2. Section 3 describes the corpora, measurement instrument, and gap analysis procedures. Section 4 presents the empirical findings. Section 5 interprets the findings and discusses limitations. Section 6 concludes with implications and directions for future research.

## 2 Literature Review

### 2.1 AI and the SDGs: What Bibliometrics Shows

This section reviews past studies on AI and the SDGs. Most ask which goals draw the most research attention.

The best-known work, such as Vinuesa et al. (2020)’s assessment of AI’s effects on the 169 SDG targets, rests on expert views or counts rather than on the texts themselves.

Broader SDG scientometric work has mapped the growth, institutional distribution, thematic structure, and interlinkages of MDG/SDG-related research across methods and time windows (Bautista-Puig et al., 2021; Cowls et al., 2021; Nedungadi et al., 2024; Singh et al., 2024). Bautista-Puig et al. (2021) chart the general SDG research landscape across higher-education and research institutions from 2000 to 2017 using a citation-based corpus of roughly 25,000 publications, finding SDG 3 (Good Health) by far the most addressed goal and SDG 7 and SDG 12 among the least. The AI-specific inventories converge on SDG 3 as dominant: Singh et al. (2024) map 20,511 AI publications (2001–2020) and rank SDG 3 and SDG 7 as the most applications-rich goals with SDG 17 the sparsest, while Cowls et al. (2021) survey 108 AI-for-social-good projects and report SDG 3 as the leader and SDGs 5, 16, and 17 each addressed by fewer than five. Nedungadi et al. (2024) apply BERTopic modelling to 1,288 big-data and AI publications (2013–2023) and likewise find SDG 3 most prominent, with SDG 9 and SDG 7 next. Across these studies, SDG 3 is the one robustly dominant goal, while the least-covered goals differ between the AI-specific inventories and the broader bibliometric picture — but none of them measures within-SDG semantic framing.

The methodological concern raised by Armitage et al. (2020) is important: SDG publication mapping depends on interpretive choices about what counts as SDG relevance, how target concepts are translated into queries or labels, and which bibliographic database is used. The embedding-based approach used here reduces some vocabulary-level brittleness of keyword dictionaries, but it remains a distributional measurement strategy whose results are conditional on corpus construction, centroid provenance, and model choice. Later benchmarking work reaches the same caution from a broader comparison of Boolean, machine-learning, and hybrid SDG-mapping systems: Kashnitsky et al. (2024) find that no single approach performs best across all validation datasets, and that no single “golden” SDG validation set exists.

Not only are SDG mapping systems methodologically contingent; the SDGs themselves are interdependent by design. Nilsson et al. (2016) observe that “the goals depend on each other” but that no operational framework specifies how, warning that treating targets in isolation “risk[s] perverse outcomes.” Single-SDG assignment is therefore a simplifying

assumption, not a reflection of how goals operate in practice.

One SDG breaks the interdependence pattern. Le Blanc (2015), in a network analysis of the proposed SDGs, treats SDG 17 (Partnerships for the Goals) as the dedicated goal for cross-cutting means of implementation — finance, technology transfer, capacity-building, and partnerships — and excludes it from cross-goal network mapping because its targets are implementation infrastructure rather than substantive domains: a distinction unique among the 17 goals. The Griggs et al. (2017) guide likewise describes SDG 17 as the main enabler of the whole framework. Unlike every other SDG, whose discourse centres on a shared substantive topic, SDG 17 policy discourse concerns coordinating and resourcing implementation itself — a conceptual space that technical AI research abstracts are not built to occupy.

Past work shows which SDGs attract AI and sustainability research, and how much that mapping depends on method and corpus. Whether research and policy texts for the same goals actually align is the subject of the next section.

## 2.2 A Three-Level Framework of Research-Policy Alignment

This paper proposes that research-policy alignment be understood as a three-level construct rather than a binary property of shared labels:


**Level 1 — lexical or bibliometric correspondence.** The dominant mode of existing SDG-mapping work, assigning texts to goals through keyword dictionaries, Boolean queries, or citation and mention counts (Armitage et al., 2020; Kashnitsky et al., 2024). Necessary but weak: shared vocabulary does not guarantee shared meaning (Bornmann et al., 2016), and keyword methods are brittle to vocabulary and database choices (Armitage et al., 2020).

**Level 2 — observed textual-semantic proximity in embedding space.** Sentence embeddings locate texts in a shared semantic space where proximity reflects distributional patterns of co-occurrence (Harris, 1954; Reimers & Gurevych, 2019), and encoder-based SDG association has been benchmarked as a competitive baseline (Gjorgjevikj et al., 2025). This level detects divergence that keyword methods miss, while remaining a proxy rather than a definitive conclusion: embedding distance mixes modality, abstraction, institutional stance, and theme (Biber, 1988), and the research-policy contrast is primarily a register contrast (Biber & Conrad, 2009). The present study is positioned at this level.

**Level 3 — substantive comparison of implied problems, actors, and interventions.** The ultimate object of alignment is whether the two discourses address the same problems, involve the same actors, and propose the same interventions (Cash et al., 2003; Guston, 2001; Haas, 1992). Because this requires interpreting causal beliefs, legitimacy,

and institutional mechanism (Caplan, 1979; Haas, 1992), it lies beyond unsupervised text analysis and is left to qualitative follow-up.

Advancing from Level 1 to Level 2 is the study’s concrete methodological contribution; Level 3 names the boundary it deliberately does not cross. Figure 1 illustrates this decomposition.



**Figure 1.** Conceptual framework: research–policy divergence separated into coverage and semantic framing dimensions, positioned within the three-level alignment construct.

Figure 3 summarises the complete data pipeline (Section 3.11).

## 2.3 Research-Policy Alignment in AI and Sustainability

This section asks whether AI research and sustainability policy line up, and why past work says they often do not.

A growing literature documents structural tensions between what AI research produces and what sustainability policy requires. Strauss et al. (2025) show that corporate AI research concentrates on pre-deployment technical problems, creating a gap between research priorities and the real-world deployment impacts that governance and policy must address. Sioumalas-Christodoulou and Tympas (2025) find that global AI metrics focus on technical and economic performance while policy discourse emphasises ethical and social dimensions aligned with SDG-related concerns. Toney-Wails et al. (2024), in the trustworthy-AI governance literature, provide a close methodological precedent, showing that policy and research communities can use shared trustworthy-AI terminology while differing in term salience, definitions, and focal concerns.

McCarthy et al. (2026) report a parallel finding in conservation science, where a multi-target SciBERT classifier applied to Ross Sea MPA research reveals that 59% of policy-prioritised research topics appear in fewer than 15 papers, demonstrating that research-policy coverage gaps are observable across domains.

Adjacent scientometric work has also used policy documents to measure research-policy connection, but at the level of mentions rather than semantic framing. Bornmann et al. (2016) examine climate-change publications in Altmetric-tracked policy documents and find that only 1.2% of 191,276 publications had at least one policy mention. Their study

is useful here as a boundary marker: policy-document mentions can indicate one narrow form of policy uptake, but they do not show how research is discussed, interpreted, or framed inside policy discourse. The present study therefore addresses a different layer of the research-policy relationship: not whether policy documents cite particular papers, but whether research and policy-facing texts assigned to the same SDG occupy nearby regions in semantic space.

A related research-evaluation literature treats the research-policy relationship through the lens of societal impact measurement. Bornmann (2017) argues that measuring societal impact is harder than measuring scientific impact because the target of impact must first be specified: research may affect policymakers, business, culture, public health, or other social domains. Policy documents are therefore useful as one possible target-oriented source for measuring impact on policymaking, but this is a different task from semantic framing comparison. The present study does not ask whether research has achieved policy impact; it asks whether research and policy-facing texts occupy similar semantic regions when organised through shared SDG labels.

These findings are consistent with theoretical accounts of research-policy misalignment. Heink et al. (2015) identify credibility, relevance, and legitimacy as determinants of whether scientific knowledge informs policy; topical relevance alone is insufficient — relevance also depends on whether research results can shape actual decisions. In sustainability science, Cash et al. (2003) make a closely related argument: knowledge systems are more likely to support action when they produce information that is simultaneously salient, credible, and legitimate, and when they manage the boundary between knowledge and action through communication, translation, and mediation. The epistemic communities framework (Haas, 1992) provides a theoretical reason why shared framing matters for policy coordination under uncertainty: expert communities can influence policy by articulating cause–effect relationships, framing issues for collective debate, proposing policies, and identifying salient points for negotiation. However, Haas’s concept is broader than semantic proximity alone, involving shared normative beliefs, causal beliefs, validity criteria, and a common policy enterprise. The present study therefore does not operationalise an epistemic community directly; it measures the weaker condition of observed semantic proximity between research and policy-facing discourse.

Cross (2013) strengthens this caution by arguing that epistemic communities should not be treated as simply present or absent: their influence varies with internal cohesion, professionalism, access to decision-makers, competition with other actors, and broader political context. This means that semantic proximity can at most indicate a possible precondition for knowledge transfer, not the existence, strength, or influence of an epistemic community.



The SDG framework is itself a politically negotiated instrument with documented limitations. Fukuda-Parr (2016) argues that global goals can have “distorting effects, redefining the meaning of development, and shaping policy by creating incentives to over-focus on target achievement.” Bexell and Jönsson (2017) find that SDG summit documents remain “state-centric with great room for state sovereignty,” avoid assigning causal responsibility for structural problems, and rely on “weak and unsystematic accountability.”

A more fundamental tension exists within the framework itself. Hickel (2019) demonstrates that “Goal 8 violates the sustainability objectives of the SDGs” — 3% annual global growth is empirically incompatible with the resource-use and emissions targets in Goals 6, 12, 13, 14, and 15. The SDGs “lack theoretical grounding” and embody a contradiction between growth and ecological sustainability.

Heilinger et al. (2024) warns that the “sustainable AI” label can be used for greenwashing where technologies are framed as contributing to the SDGs without sufficient basis. Heras-Saizarbitoria et al. (2022), analysing 1,370 sustainability reports, find that most organisations’ SDG engagement is superficial — “the SDGs only serve to add colour and fancy icons to reports” — suggesting that SDG-washing is a measurable empirical pattern. This supports the caution that reported alignment should not be conflated with substantive commitment.

Measuring semantic alignment with SDG centroids therefore uses a benchmark that is itself politically negotiated and internally inconsistent — not a neutral standard. Gibbons et al. (1994) distinguish Mode 1 knowledge production, where problems are set and solved within academic disciplinary interests, from Mode 2, where problems arise in contexts of application; to the extent that research follows Mode 1, its agendas may reflect internal disciplinary dynamics rather than external policy needs. This observation also operates alongside a further ambiguity: what it means to be “aligned” with a framework whose component goals are in tension with each other, and it suggests that divergence is not uniformly problematic — where research explores possibilities that policy discourse has not yet absorbed, distance may reflect productive exploration rather than failure of alignment.

The two-communities theory (Caplan, 1979) holds that researchers and policy makers live in separate worlds with different values, reward systems, and languages, so divergence is the default state of research-policy interaction rather than a failure. Taken together, this literature indicates that the divergence between AI research and sustainability policy is a substantive feature of the two knowledge systems rather than an artefact of measurement. Existing studies document the gap through coverage, citation, and impact-based methods, while theory explains it through legitimacy, expert communities, and the structural divide between academic and policy communities. None of these studies measures semantic framing within a single SDG. This study measures how close research and policy texts lie

in semantic space; that closeness is a first step, not proof of alignment.

## 2.4 Semantic Methods: Rationale and Cautions

This section explains the method behind this study and why it needs care.

The use of Sentence-BERT (Reimers & Gurevych, 2019) and related transformer-based encoders for semantic comparison has been validated across several domains relevant to this study, but applying them to compare two distinct discourse corpora — research abstracts and policy documents — is a novel usage rather than a directly benchmarked one. This method follows the distributional hypothesis that meaning can be described through patterns of linguistic co-occurrence, a principle whose modern computational implementations trace back to Harris (1954)’s foundational observation that the distribution of a linguistic element is “the sum of all its environments.” Sentence-BERT and related transformer-based encoders are therefore used here as pragmatic tools for large-scale semantic comparison, while the interpretation of distance remains descriptive rather than causal. Gjorgjevikj et al. (2025) benchmarked sentence encoders specifically for SDG association tasks, finding that general-purpose models from the sentence-transformers family perform competitively as baselines, and that domain-specific fine-tuning further improves predictive performance while reducing sensitivity to description length.

Rodriguez et al. (2023)’s conText framework provides a closer computational-social-science precedent for comparing meaning across contexts. Their embedding-regression approach estimates how focal terms are used differently across covariates such as party, time, or corpus, while explicitly treating embedding-based “meaning” as a thin distributional description rather than a causal model of human interpretation. This study does not adopt their word-level regression estimator, but it follows the same broad premise that semantic representations can be used descriptively to compare how different discourse communities frame shared terms or categories.

McCarthy et al. (2026) supply a yet closer empirical precedent for the present design. Applying a multi-target SciBERT classifier to Ross Sea MPA research and policy-prioritised topics, they find that 59% of policy-prioritised research topics appear in fewer than 15 papers, showing that research–policy coverage gaps are measurable with transformer encoders outside the AI–SDG setting. This study extends that logic from a single conservation domain to the full SDG landscape, and from coverage to within-SDG semantic framing.

An older corpus-linguistic tradition gives a further reason for caution when interpreting cross-corpus semantic distance. Biber’s multidimensional analysis of spoken and written English showed that register differences are not simple binary contrasts, but continuous patterns of co-occurring linguistic features across genres (Biber, 1988). For the present

study, this matters because research abstracts and policy documents may differ not only in substantive topic, but also in abstraction, modality, institutional stance, and rhetorical function. Embedding distance is therefore treated as a measure of observed textual-semantic proximity, not as a clean measure of topic alone.

Embedding models let us compare research and policy texts, but applying them across two corpora is new. They represent meaning through distribution, not cause, and cannot separate topic from register. The measures used here are descriptive: they report how close or distant the two corpora are in semantic space, not whether that closeness is real alignment.

## 2.5 Research Gap

To the best of my knowledge, the literature reviewed above has not produced: (i) a systematic comparison of how academic AI-for-sustainability research and institutional SDG policy discourse differ in SDG coverage and within-SDG semantic framing — prior work measures research coverage, citation patterns, or policy-document uptake, but does not compare the two corpora directly on these dimensions; or (ii) an examination of whether coverage gaps and semantic gaps are positively related, weakly related, or inversely related. These are the contributions of the present study, and they answer the central research question set out in the Introduction. Together, they provide a systematic framework for distinguishing research-policy coverage from research-policy semantic proximity within a shared SDG representation. This framework is deliberately diagnostic: it maps where observed divergence is visible, without classifying whether any given gap reflects productive or problematic disconnection — the latter would require qualitative follow-up beyond the scope of embedding analysis.

# 3 Methodology

## 3.1 Research Corpus

This study uses only publicly available text data and codebase. It does not involve human participants or personal data. No ethical approval was required.

The research corpus consists of 3,105,144 English-language research abstracts retrieved from the OpenAlex API (Priem et al., 2022), covering publications between 2018 and 2025. Abstracts were retrieved by combining OpenAlex’s native SDG work filter with four AI terms (`artificial intelligence`, `machine learning`, `deep learning`, `neural network`) — 68 queries (17 SDGs  $\times$  4 terms). Complete query parameters are documented in `code/0_fetch/fetch_openalex.py`. OpenAlex SDG metadata defines the candidate universe; the SDG assignments used in the analysis are produced by the independently

validated supervised classifier (Section 3.4).

To retrieve the AI research corpus, we deployed a high-precision four-term query string, namely **artificial intelligence, machine learning, deep learning, and neural network**. This term-based retrieval follows the established keyword-based practice for defining an AI publication corpus. Elsevier’s data-driven characterisation of AI research, for instance, draws on over 700 field-specific keywords across some 600,000 documents (Hellwig et al., 2019). We acknowledge that comprehensive AI bibliometric mapping instead favours recall. The OECD.AI Observatory methodology (OECD, 2024) defines a paper as AI when it carries the Artificial Intelligence or Machine Learning field of study in the Microsoft Academic Graph (MAG), a scholarly knowledge graph whose Fields of Study taxonomy OpenAlex later inherited and extended by mapping it to Wikidata identifiers (Priem et al., 2022; Wang et al., 2020). This concept-tagging schema underpins the OECD.AI data pipeline for identifying AI publications (OECD, 2024). OpenAlex (Priem et al., 2022) is the open-access successor to MAG, so our concept retrieval reuses the same field-of-study definition via OpenAlex concept tags for Artificial Intelligence and Machine Learning. Because our objective is macro-level SDG semantic alignment rather than exhaustive publication counting, we prioritise precision at the retrieval stage to keep non-AI noise out of the embedding space. The corpus is therefore conditional on this term selection, and retrieval boundary sensitivity is a scope condition that future work with broader taxonomies may revisit.

A direct test of this scope condition re-ran retrieval using field-of-study membership for the same two AI fields of study (OECD, 2024), then re-scored the resulting 50,000-paper corpus through the identical pipeline. This 50,000-paper scale is justified on two independent grounds. First, it sits on the empirical variance plateau established by our sample-stability ladder (Table 9): at 50,000 papers the coverage-profile standard deviation is 0.001 and the mean within-SDG semantic gap is 0.358, within 0.002 of the full-corpus estimate of 0.356 (full corpus  $N = 3,105,144$ ), while policy-text calibration bias has also converged (0.113). Second, 50,000 papers exceeds every corpus size used in prior SDG classification and mapping benchmarks, namely the OSDG community dataset (30,534 texts; Pukelis et al. (2022)), the SDG benchmark encoder evaluation corpus (35,001 papers; Gjorgjevikj et al. (2025)), the SDGi corpus (20,267 segments; SDGi (2024)), the Aurora expert-validated set (4,971 texts; Vanderfeesten et al. (2020)), and the 6,000-paper Kashnitsky et al. benchmark (Kashnitsky et al. (2024)), while remaining far smaller than the full 3.1-million-paper corpus and thus a computationally tractable substrate for OpenAlex concept tagging.

The within-SDG coverage-gap and semantic-gap rankings of the concept corpus correlate strongly with the keyword-retrieved baseline (SDG-share Spearman  $\rho = 0.936$ , coverage-gap rank  $\rho = 0.96$ , semantic-gap rank  $\rho = 0.89$ ; Appendix G). This shows the reported divergence patterns are not an artefact of the four-term query choice. Schoeberl et al. (2023)

independently reach the same methodological conclusion for AI-research identification, namely that results are sensitive to the retrieval method and that supervised classification is preferred, and the present LR scoring is exactly that classifier-based step applied to SDG assignment.

The 2018 start date reflects the post-AlexNet (Krizhevsky et al., 2012) emergence of deep learning as the dominant paradigm and a reasonable research-uptake lag after the SDGs’ adoption in 2015.

An important asymmetry follows: the research corpus is a tight intersection of AI and SDG content ( $AI \cap SDG$ ), while the policy corpus (Section 3.2) is a broader union covering institutional SDG discourse and AI governance. Between-SDG rankings should be interpreted as conditional on this asymmetric construction. Following Armitage et al. (2020), this retrieval universe is a method-dependent corpus boundary, not a neutral population of all SDG-relevant AI research. After deduplication and removal of records lacking usable abstracts, 3,105,144 abstract-derived text segments were retained. No balancing by SDG or citation count is applied — trimming would distort the coverage profile the study aims to measure.

### 3.2 Policy Corpus

The policy corpus contains 40,547 text segments from three collections, representing 2,455 unique source documents. It should be read as mixed *international institutional SDG policy discourse* rather than the full universe of AI sustainability policy:

1. **Curated AI and SDG policy documents** (13,143 segments, 94 documents): UN, OECD, EU, UK, and AU AI governance frameworks; IPCC assessment reports; SDSN sustainable development reports; and national AI strategies. These are AI-governance instruments by construction, forming the innovation-facing slice of the corpus.
2. **SDGi Voluntary National Review and Voluntary Local Review corpus** (SDGi, 2024) (21,346 segments, 313 documents): National and local self-assessments submitted to the UN High-Level Political Forum. The creators describe it as the most comprehensive multilingual collection of SDG-labelled texts to date, spanning all 17 goals and centred on SDG implementation rather than any single sector.
3. **UN General Debate Corpus** (Harvard Dataverse, 2024) (6,128 segments, 2,048 documents): SDG-relevant passages from General Debate speeches, sessions 70–80 (2015–2024). Its creators analyse SDG coverage directly and report climate action (SDG 13), peace and governance (SDG 16), and partnerships (SDG 17) among the most prominent themes — a climate–governance–partnership orientation the analysis below also observes.

Each source document was segmented using the transformer model’s own tokenizer: NLTK sentence-boundary detection (`sent_tokenize`) is followed by greedy accumulation of sentences up to a token budget of `max_seq_length - 10` (margin 10, empirically verified to reduce truncation to near zero). Segments falling below 20 words (`MIN_SEGMENT_WORDS`) are discarded; a single sentence exceeding the token budget is kept whole because mid-sentence splitting is not permissible. These raw collection totals (13,143, 21,346, 6,128) reflect post-segmentation counts before cross-source deduplication. For the UN General Debate Corpus, paragraph-level keyword filtering precedes segmentation, so only SDG-relevant speech passages enter the pipeline — this pre-filtering is why UNGDC’s raw segment count is much smaller than its source-speech count. After deduplication by exact text match (SDGi retained first, removing 70 curated-source segments that duplicated SDGi text), 40,547 unique segments remained. Segment identifiers encode source document identity for document-level weighting.

**Source heterogeneity note.** The three collections differ in genre and scope — the curated subset is AI-governance-focused, while SDGi and UNGDC reflect broader institutional SDG discourse. The full-corpus policy distribution should be read accordingly; the curated sub-corpus (94 documents) is the more relevant reference for AI-governance-specific interpretation. Appendix H.2 reports a source-family sensitivity check.

### 3.3 Embedding Model and Normalisation

All texts were encoded using `all-mpnet-base-v2` (Sentence-Transformers, 2021) from the sentence-transformers library (Reimers & Gurevych, 2019), producing 768-dimensional dense vector representations. This pre-trained model achieves competitive performance on semantic textual similarity benchmarks and requires no fine-tuning on the data under analysis — the embedding space is a stable pre-existing coordinate system. All embeddings were L2-normalised to unit vectors so that cosine similarity between any two vectors equals their dot product.

During pipeline development, a truncation issue was discovered and corrected: texts exceeding the encoder’s maximum sequence length were being silently truncated before embedding, discarding tail content. The fix uses the model’s own tokenizer during segmentation (Section 3.2) to enforce a token budget of `max_seq_length - 10`, verified by the `verify_truncation_rate` utility to produce near-zero post-segmentation truncation. The full corpus was re-embedded after the fix; all reported results use post-fix embeddings.

### 3.4 Supervised Reference Classifier

The canonical instrument is a supervised multinomial logistic regression (LR) classifier trained on pooled single-label reference data from five independent annotated corpora:

- **OSDG Community Dataset** (Pukelis et al., 2022) (30,534 texts, SDGs 1–16): a crowdsourced citizen-science collection of excerpts from reports, policy documents, and abstracts, each validated by multiple volunteers.
- **SDG Classification Benchmark** (SDG Classification Expert Group, 2025) (616 expert-verified texts, SDGs 1–17): short text snippets independently annotated by multiple domain experts with consensus resolution.
- **IISD SDG Knowledge Hub** (Wulff et al., 2023) (6,125 texts, SDGs 1–17): IISD journalism and policy commentary on 2030 Agenda implementation.
- **SDGi Corpus** (SDGi, 2024) (20,267 texts, SDGs 1–17): government Voluntary National and Local Reviews submitted to the UN.
- **Aurora dataset** (Vanderfeesten et al., 2020) (4,971 expert-validated texts, SDGs 1–17): academic abstracts validated by domain researchers.

All five corpora supply human- or expert-validated labels rather than machine-generated ones, and they differ in provenance. OSDG and SDGi are policy-adjacent: OSDG volunteers confirm a suggested label across multiple annotators, and SDGi compiles government Voluntary National and Local Reviews submitted to the UN. The SDG Knowledge Hub is journalism and reporting; Aurora is academic, expert-validated text. The mix broadens the label base, but OSDG alone contributes about 49% of all labelled texts and draws on UN-adjacent documents, so the reference data remain weighted toward policy vocabulary. No high-quality graded relevance sets exist for SDG classification, making this combined, multi-source collection the best available resource (Kashnitsky et al., 2024).

The pooled reference data (n=62,513 after preprocessing) were split into training and held-out test sets (n=52,779 train, n=9,734 test) by stratified document-grouped splitting: all texts from the same source document are kept in the same fold, enforcing strict separation between training and evaluation. The champion LR (C=10.0, L2 penalty, class\_weight=None, solver=lbgfs) was selected by 5-fold GroupKFold cross-validation on the training pool; the full grid-search protocol, hyperparameter rationale, exclusion of alternative model families (RF, XGB), and the comparison with the MLP robustness check are reported in Appendix D. The main text states plainly that LR is the canonical method; the MLP is a robustness check, not an alternative champion.

The validation task evaluates whether the classifier assigns held-out test examples to the correct SDG; it does not evaluate a rejection option for non-SDG texts. This is acceptable because the research corpus is first restricted to an OpenAlex AI-and-SDG universe — the classifier then answers a relative question: among the 17 SDGs, which is this text closest to?

### 3.5 Scoring and Assignment

A text is assigned to the SDG with the highest predicted probability from the trained LR classifier (argmax over 17 classes). Every assignment is traceable to 768 signed logistic-regression coefficients per SDG class — a transparent decision boundary with no hidden layers, no dropout, and no random seed sensitivity. Macro-F1, the unweighted mean of per-SDG scores, is the primary validation metric; the scores are reported in Section 4.1 alongside SDG-specific concerns that inform the results below.

LR argmax assigns each text to the SDG whose decision boundary it falls on the positive side of. The method is simple, fast, and avoids the hidden-layer opacity of non-linear classifiers. I chose it deliberately: the study needs a transparent tool to compare corpora, not a classifier built to label single documents. LR additionally provides calibrated per-class probability outputs, which can serve as a soft-assignment comparison against the hard argmax assignment; the canonical analysis uses hard argmax for direct interpretability.

Additional diagnostic checks — within-corpus diagnostics, policy source-family sensitivity, and an SDG 4 lexical audit — are reported in Appendix B. The canonical analysis uses hard LR argmax assignment as a transparent and reproducible anchoring procedure.

### 3.6 Classifier Calibration Across Discourse Types


Because the classifier is trained on pooled reference data that are themselves policy-adjacent (Section 3.4), the scoring instrument behaves differently on the two corpora. Policy segments receive systematically higher LR predicted probabilities than research abstracts (mean 0.768 vs 0.656, a 0.113 gap). The gap has two causes the data cannot separate: a genre artefact (the reference data skew toward policy vocabulary) or a substantive difference (policy texts address SDG implementation more directly). Either way the asymmetry does not threaten the findings: hard argmax assignment uses only within-corpus ranking, not absolute probability levels, so the directional gap does not reorder goals. The asymmetry is not an SDGi circularity effect: the training set and the scored policy corpus are disjoint by construction.

The combined PCA landscape (Figure 2) visualises this asymmetry: the policy cloud hugs the 17 reference-classifier centroids more tightly than the research cloud. The projection is descriptive only (the first two components explain 9.2% and 3.7% of variance); all reported cosine distances remain in the original 768-dimensional space.

### 3.7 Coverage Gap Analysis

For each corpus item, the predicted SDG is the LR argmax (hard assignment, Section 3.5). The coverage profile is the proportion of items assigned to each SDG.





**Figure 2.** PCA semantic landscape of the combined research-policy embedding space. Descriptive two-dimensional projection only; formal semantic distances are computed in the original embedding space.

The policy corpus is not a flat set of documents: its 40,547 segments are grouped into 2,455 source documents, and those documents vary enormously in length (SDGi: 21,346 segments across 313 documents; curated AI/SDG: 13,073 across 94; UNGDC: 6,128 across 2,048). A segment-weighted profile would let a handful of long reports dominate the count, so the analysis instead computes a *document-weighted* profile in which every source document counts equally regardless of length. Concretely: (1) the segment score vectors of each document are averaged into one document-level vector; (2) that vector is hard-assigned to its top SDG; (3) the coverage profile is the share of documents — not segments — assigned to each SDG. The research corpus uses the same logic at the paper level, where each of the 3,105,144 abstract texts is already a single unit. Document-weighted profiles are the canonical representation; segment-weighted profiles are retained as a diagnostic comparison.

The coverage gap for SDG  $j$  is defined as:

$$\text{CoverageGap}_j = \left| \text{ResearchCoverage}_j - \text{PolicyCoverage}_j^{\text{docwt}} \right|$$

The signed version of the same quantity is:

$$\text{SignedCoverageGap}_j = \text{ResearchCoverage}_j - \text{PolicyCoverage}_j^{\text{docwt}}$$

A positive value means research dominates SDG  $j$ ; a negative value means policy dominates. These two definitions, together with research coverage and policy coverage, are the four predictors reported in Table 2.

### 3.8 Semantic Gap Analysis

The semantic gap measures within-SDG divergence: given that both corpora have texts assigned to SDG  $j$ , how differently do they discuss it? For each SDG  $j$ , the analysis collects all assigned research papers and all assigned policy segments (capped at 50 per source document to prevent single-document dominance — SDSN contributes up to 3,179 segments otherwise), computes the L2-normalised mean embedding of each cluster ( $\mathbf{c}_j^{\text{res}}$ ,  $\mathbf{c}_j^{\text{pol}}$ ), and defines:

$$\text{SemanticGap}_j = 1 - (\mathbf{c}_j^{\text{res}} \cdot \mathbf{c}_j^{\text{pol}})$$

A gap of 0 indicates perfectly overlapping semantic directions; a gap approaching 1 indicates orthogonal framing.

The policy corpus is represented as short segments rather than whole documents, so a per-SDG policy sub-centroid is a mean over all assigned segments. Without limitation, a handful of very long documents would dominate the cluster: within the curated AI/SDG collection, SDSN sustainable development reports alone contribute up to 3,179 segments, and the SDGi Voluntary National and Local Review reports together span 21,346. I therefore apply a *per-document segment cap*: for each source document, at most  $K$  of its segments are randomly sampled (seeded, hence reproducible) before the sub-centroid is built, so that no single document can outweigh the rest. Random rather than positional sampling (e.g. taking the first  $K$  segments) avoids biasing the centroid toward a document’s opening or closing passages. The canonical cap is  $K = 50$  ( $\approx 7,500$  words), a conservative choice well above the median of roughly 14 segments per document. The cap choice was checked by repeating the computation at an alternative cap of 20 segments per document and with no cap at all; the SDG gap rankings are near-identical to the canonical cap-50 results (rank correlation  $\rho = 0.968$ ; Table 3, Section 4.5), so the findings are robust whether or not the cap is applied.

The semantic gap is not a clean measure of substantive disagreement. Following Biber and Conrad’s distinction between register, genre, and style (Biber & Conrad, 2009), I treat the research–policy contrast primarily as a register contrast: the two corpora differ in situation

of use, audience, and communicative purpose, not merely in topic. Embedding distance therefore mixes a genuine within-SDG substantive divergence (the two corpora discuss the same SDG differently) with register divergence (research abstracts and policy documents differ in abstraction, modality, and institutional stance) (Biber, 1988). The raw centroid distance is retained as the primary measure because it directly corresponds to this observed framing difference. Separating the substantive component from the register component is an open problem. Appendix F reports a register-adjusted sensitivity analysis: learning one global research-policy direction in embedding space and subtracting its projection from both corpora before recomputing the gaps. The adjustment compresses the gaps moderately but preserves the relative ordering of high-gap and low-gap SDGs. These adjusted estimates bound the register-sensitive component but do not replace the raw gap.

### 3.9 Coverage–Semantic Interaction

The first hypothesis (H1) decomposes into four named sub-hypotheses, each correlating a distinct coverage predictor with the within-SDG semantic gap across the 17 SDGs: H1a, coverage gap  $\leftrightarrow$  semantic gap; H1b, research-policy dominance  $\leftrightarrow$  semantic gap; H1c, research coverage  $\leftrightarrow$  semantic gap; and H1d, policy coverage  $\leftrightarrow$  semantic gap. H1a is the primary, motivating hypothesis: the absolute research-policy coverage gap is expected to predict semantic divergence. H1b is its directional complement, testing whether the signed balance of research and policy attention (which side dominates a goal) also predicts divergence. I compute Pearson  $r$  and Spearman  $\rho$  between each coverage predictor and semantic gap scores across all 17 SDGs, with sensitivity analyses excluding SDG 4 (see Section 3.10). Given  $n = 17$ , the tests have limited power to detect anything but a large linear effect; this is a scope condition.

### 3.10 Key Assumptions and Mitigations

This subsection records the design-level assumptions that shape the comparison before any measurement takes place. Two such choices are methodological decisions rather than empirical findings: how text units are handled across the two corpora, and how a known limitation of the SDG 4 reference labels is treated.


**Text-unit handling is a design choice, not a result.** Research papers are single abstract-length texts (mean  $\sim 200$  words); policy items are 150–300-word segments from documents of highly variable length. Document-weighting and per-document segment caps are adopted up front to keep the two corpora comparable; they are fixed procedures, not findings derived from the data.

**Labelled-data limitation on SDG 4.** The research coverage assigns 14.6% of papers to SDG 4 (Quality Education, see later Section 4.2), flagged as artefact-affected because ML

papers share core vocabulary (“learning”, “training”, “model”) with the OSDG education corpus — a case of lexical brittleness in keyword-based semantic mapping (Armitage et al., 2020). Appendix A.1 confirms that SDG 4-assigned texts contain education-specific terms more often than matched non-SDG4 samples, but co-occur with ML-learning vocabulary. The result is retained as a learning-related SDG 4 assignment rather than a clean education estimate. Sensitivity analyses excluding SDG 4 are provided for all correlation results. This is a limitation of the reference labels, carried through to the results and mitigated there by artefact-flagged sensitivity analyses, not resolved at the design stage.

### 3.11 Methodology Summary

Figure 3 summarises the end-to-end methodology in three parallel data tracks. Each track begins with source-specific retrieval, then passes through preprocessing (cleaning, English filtering, 20-word minimum) and token-aware segmentation (`segment_text`, margin=10, min 20 words). In Track 1 (Reference/Training), three of the five corpora (KH, SDGi-labels, Aurora) undergo segmentation while OSDG and Benchmark bypass this step as they already consist of short single-label texts, yielding a pooled 62,513 texts. These are split via per-source stratified GroupKFold (52,779 train, 9,734 test) and used to train the LR classifier ( $C = 10$ , L2, macro- $F_1 = 0.823$ ). Track 2 (Research) segments all 3,105,144 abstract-derived passages. Track 3 (Policy) segments each of the four policy sources (scrape, manual, UNGDC, SDGi) individually and then merges them with cross-source deduplication (SDGi-first), yielding 40,547 unique policy segments. All three tracks are embedded with `all-mpnet-base-v2` (768-D, L2-normalised). The trained classifier scores both the Research and Policy embeddings via argmax over 17 SDG classes. Detailed per-corpus counts are reported in Sections 3.1, 3.2, and 3.4.



**Figure 3.** End-to-end methodology pipeline in three parallel data tracks. All tracks pass through preprocessing and token-aware segmentation; OSDG and Benchmark in Track 1 bypass segmentation (short single-label texts). Policy segments four sources then merges (40,547). Research segments all abstracts (3,105,144). The trained LR classifier scores both Research and Policy embeddings for downstream gap analysis.

## 4 Results

This section reports what the instrument measured: the validation scores, the coverage and semantic gaps, the H1 correlation tests, and the robustness of those rankings. It describes the findings; it does not interpret them. Interpretation, limits, and implications follow in the Discussion (Section 5).

### 4.1 Supervised Reference Classifier Validation

The LR classifier was evaluated on the held-out test set ( $n=9,734$ , document-grouped to prevent source leakage). It reaches macro-F1 = 0.823 and micro-F1 = 0.8369, far above the 1/17 random baseline (a 17-class random guess scores  $1/17 \approx 0.059$ ). Per-class F1 values from the test set are reported in Table 1.

**Table 1.** LR test-set F1 scores ( $n=9,734$ ).

| SDG                           | LR test F1   |
|-------------------------------|--------------|
| SDG 1 (No Poverty)            | 0.758        |
| SDG 2 (Zero Hunger)           | 0.829        |
| SDG 3 (Good Health)           | 0.910        |
| SDG 4 (Quality Education)     | 0.889        |
| SDG 5 (Gender Equality)       | 0.883        |
| SDG 6 (Clean Water)           | 0.878        |
| SDG 7 (Clean Energy)          | 0.874        |
| SDG 8 (Decent Work)           | 0.678        |
| SDG 9 (Industry & Infra.)     | 0.780        |
| SDG 10 (Reduced Inequalities) | 0.637        |
| SDG 11 (Sustainable Cities)   | 0.789        |
| SDG 12 (Responsible Cons.)    | 0.786        |
| SDG 13 (Climate Action)       | 0.843        |
| SDG 14 (Life Below Water)     | 0.906        |
| SDG 15 (Life on Land)         | 0.870        |
| SDG 16 (Peace & Justice)      | 0.886        |
| SDG 17 (Partnerships)         | 0.797        |
| Macro-F1 (SDGs 1–17)          | <b>0.823</b> |

*Notes:* LR test F1 on the held-out pooled test set (all reference sources). The table shows the canonical 17-way classifier used throughout the main analysis.


The LR achieves consistently strong per-SDG performance: SDG 3 ( $F1 = 0.910$ ), SDG 14 (0.906), and SDG 4 (0.889) are the most confidently classified (each above 0.89). The weakest classes are SDG 10 (Reduced Inequalities,  $F1 = 0.637$ ) and SDG 8 (Decent Work,  $F1 = 0.678$ ), consistent with their known semantic breadth and the overlapping cluster structure of adjacent goals (see centroid-similarity heatmap, Figure 7). Most goals score between 0.76 and 0.89, indicating reliable discrimination across the 17-class space. SDG 17 ( $F1 = 0.797$ ) is no longer the weakest link — the supervised decision boundary separates partnership discourse more cleanly than the centroid average did — confirming that the goal’s broad mandate resists compact semantic definition regardless of method.

The MLP robustness check (Appendix D) achieves a comparable test macro-F1 (MLP: 0.832; LR: 0.823), confirming that architecture choice is not the binding constraint on classification performance. The gap between LR and MLP is modest and does not justify the complexity trade-off — LR provides 768 signed coefficients per class, each directly attributable to a semantic dimension, while MLP’s hidden layers diffuse this information across non-linear transformations.

**In summary**, the LR classifier separates the 17 SDGs with consistent and interpretable performance. The weakest classes — SDG 10 and SDG 8 — are known to sit in dense centroid clusters, and SDG 17, while no longer the lowest-F1 goal in the LR test set, remains the most method-sensitive SDG in the cross-sensitivity analysis. These scores mark a ceiling for the embedding-plus-linear-classifier approach; MLP results confirm that the ceiling is not an artefact of linear separability assumptions.

## 4.2 Coverage Gap: What SDGs Each Corpus Prioritises

Figure 4 presents the document-weighted policy coverage profile alongside the research coverage profile. The two profiles are markedly different.



**Figure 4.** Coverage profiles by SDG.

**Policy corpus.** Three SDGs dominate the document-weighted profile of the full institutional policy corpus: SDG 16 (Peace, Justice and Strong Institutions, 35.0%), SDG 13

(Climate Action, 18.8%), and SDG 17 (Partnerships for the Goals, 23.3%), accounting for 77.1% of policy documents. This ordering is consistent with the literature review’s account of SDG 17 as the cross-cutting means-of-implementation goal whose discourse concerns coordination and resourcing rather than a substantive topic (Griggs et al., 2017; Le Blanc, 2015); the climate–governance emphasis also matches the SDGi and UNGDC policy components described in Section 3.2. The SDG 16 lead is notable: under the supervised classifier, governance and justice discourse emerges as the most prevalent policy theme, ahead of both climate and partnerships. Appendix H.2 shows that this full-corpus profile is source-family sensitive: the curated AI/SDG sub-corpus is much more innovation-heavy, while the broader SDGi and UNGDC components are more climate-, partnerships-, and governance-oriented.

**Research corpus.** The research profile is broader, but it still concentrates strongly on a smaller set of SDGs. The largest research shares are SDG 3 (Good Health, 25.3%), SDG 9 (Innovation, 18.3%), and SDG 4 (Quality Education, 14.6%), followed by SDG 16 and SDG 11 at 11.3% and 4.8% respectively. The SDG 3 lead corroborates the literature review, where SDG 3 is the one goal every major AI and sustainability bibliometric study finds dominant — Bautista-Puig et al. (2021), Singh et al. (2024), Cowls et al. (2021), and Nedungadi et al. (2024) all rank it first, with SDG 9 and SDG 7 typically next. The SDG 4 share is the one point of disagreement: prior inventories do not rank SDG 4 among the leaders, so its third-place finish here is the same result flagged as likely inflated by machine-learning vocabulary (Section 3.10); read cautiously, the research profile agrees with the literature on SDG 3 and SDG 9 but over-states SDG 4. Even so, the research corpus clearly does not mirror the policy-side dominance of SDGs 16, 13, and 17.

**Largest coverage mismatches.** The five largest absolute coverage gaps are SDG 16 (0.237), SDG 3 (0.231), SDG 17 (0.228), SDG 13 (0.146), and SDG 9 (0.153). SDG 16 is now the most asymmetric goal — policy discourse strongly emphasises governance and institutions while research barely touches it — followed by the familiar research-over-policy dominance in health (SDG 3).


Taken together, three observations support reading the coverage profiles as a real signal rather than a measurement artefact. First, the instrument reaches macro-F1 = 0.823 on the held-out benchmark (Section 4.1), far above the 1/17 random baseline, so assignments are not random. Second, the research-side ranking reproduced here — SDG 3 dominant, with SDG 9 and SDG 7 next — matches the ordering reported independently by four prior bibliometric studies of AI and sustainability research Bautista-Puig et al. (2021), Singh et al. (2024), Cowls et al. (2021), Nedungadi et al. (2024) (Section 2), which used keyword and citation methods rather than embeddings. Third, this is all while the classifier is trained on reference corpora that are themselves policy-adjacent (Section 3.4). To summarize, that the research-covered SDG ranking nonetheless converges with independent,



methodologically unrelated literature is the strongest available evidence that the classifier anchors a meaningful semantic coordinate system, not a policy-specific artefact.

### 4.3 Semantic Gap: How Differently Research and Policy Discuss Each SDG

Figure 5 shows the within-SDG semantic gap for all 17 SDGs. The gap ranges from 0.219 (SDG 17) to 0.478 (SDG 13), a span of 0.259 cosine units. A large gap means research and policy use different words and concepts for the same goal; a small gap means they frame it the same way. The largest gap is about 4.5 times the policy–research score gap (0.113, Section 3.6). SDG 13 has the largest gap, followed closely by SDG 3 and SDG 11. The ranking is stable when the segment cap is 20 per document or removed (rank correlation  $\rho = 0.968$ ; Table 3), so the result does not depend on the sampling parameter.



**Figure 5.** Within-SDG semantic gap by SDG, sorted descending. Dashed line = median (0.333); dotted line = mean (0.356). Dark bars exceed the median; light bars fall below it.

**Largest semantic gaps.** SDG 13 (Climate Action, 0.478), SDG 3 (Good Health, 0.464), SDG 11 (Sustainable Cities, 0.460), SDG 1 (No Poverty, 0.435), and SDG 7 (Clean Energy, 0.434) show the widest gaps. The top-gap SDGs are a mix of climate, health, and infrastructure goals where research and policy discourse diverge most sharply. This ordering is stable across the segment-cap sensitivity check (Table 3). Notably, SDG 17 (Partnerships, 0.219) now shows the smallest semantic gap — the supervised decision

boundary separates partnership discourse more cleanly than the old centroid method, moving it from the most to the least divergent goal.

**Smallest semantic gaps.** SDG 17 (0.219), SDG 15 (Life on Land, 0.229), SDG 4 (Quality Education, 0.262), SDG 6 (Clean Water, 0.282), and SDG 9 (Industry and Infrastructure, 0.289) show the narrowest gaps, meaning research and policy frame these goals in close terms. A narrow gap does not mean balanced attention: SDG 9 is still research-dominant, SDG 17 is policy-dominant, and SDG 4 carries the learning-vocabulary concern (Appendix A.1). Appendix B.1 shows example texts for high- and low-gap cases.

#### 4.4 Coverage–Semantic Interaction: Two Distinct Dimensions of Observed Divergence

Table 2 reports the H1a–H1d correlation tests between coverage measures and semantic gap scores across the SDGs.

The four predictors are research coverage, policy coverage, the absolute research–policy coverage gap, and the signed research–policy dominance. Each panel shows the result across all 17 SDGs together with tests dropping SDG 4 and SDG 17.


**Table 2.** H1a–H1d correlation tests between coverage predictors and the within-SDG semantic gap.

| Config                                                                          | Pearson $r$ | $p$   | Spearman $\rho$ | $p$   |
|---------------------------------------------------------------------------------|-------------|-------|-----------------|-------|
| <b>H1a. Coverage gap <math>\leftrightarrow</math> Semantic gap</b>              |             |       |                 |       |
| Primary observed SDGs ( $n=17$ )                                                | -0.083      | 0.751 | 0.174           | 0.504 |
| Excluding SDG 4 ( $n=16$ )                                                      | -0.046      | 0.865 | 0.268           | 0.316 |
| Excluding SDG 17 ( $n=16$ )                                                     | 0.114       | 0.673 | 0.341           | 0.196 |
| <b>H1b. Policy–research dominance <math>\leftrightarrow</math> Semantic gap</b> |             |       |                 |       |
| Primary observed SDGs ( $n=17$ )                                                | 0.153       | 0.558 | 0.078           | 0.765 |
| Excluding SDG 4 ( $n=16$ )                                                      | 0.246       | 0.358 | 0.182           | 0.499 |
| Excluding SDG 17 ( $n=16$ )                                                     | -0.060      | 0.825 | -0.082          | 0.762 |
| <b>H1c. Research coverage <math>\leftrightarrow</math> Semantic gap</b>         |             |       |                 |       |
| Primary observed SDGs ( $n=17$ )                                                | 0.085       | 0.746 | 0.250           | 0.333 |
| Excluding SDG 4 ( $n=16$ )                                                      | 0.188       | 0.486 | 0.385           | 0.141 |
| Excluding SDG 17 ( $n=16$ )                                                     | 0.006       | 0.982 | 0.100           | 0.713 |
| <b>H1d. Policy coverage <math>\leftrightarrow</math> Semantic gap</b>           |             |       |                 |       |
| Primary observed SDGs ( $n=17$ )                                                | -0.121      | 0.643 | 0.025           | 0.926 |
| Excluding SDG 4 ( $n=16$ )                                                      | -0.155      | 0.567 | 0.079           | 0.770 |
| Excluding SDG 17 ( $n=16$ )                                                     | 0.075       | 0.784 | 0.206           | 0.444 |

The evidence does not support H1a: across the 17 SDGs, coverage divergence and semantic

divergence show no detectable association ( $r = -0.083$ ,  $p = 0.751$ ;  $n = 17$ ). Given the small number of SDG-level observations, this should be interpreted as evidence against a strong monotonic relationship rather than evidence of exact independence. H1c likewise shows no detectable association: research coverage and semantic gap are uncorrelated ( $r = 0.085$ ,  $p = 0.746$ ). Two exploratory tests — H1d, policy coverage ( $r = -0.121$ ,  $p = 0.643$ ), and H1b, signed dominance ( $r = 0.153$ ,  $p = 0.558$ ) — are also near-zero. All correlations remain small and non-significant under both exclusions; however, H1a and H1b flip sign when SDG 17 is dropped (Table 2), consistent with SDG 17’s broader sensitivity to measurement configuration (§4.5).

Figure 6 shows each SDG. The most striking pattern is the absence of a pattern: SDGs with large coverage gaps span the full range of semantic gaps (e.g. SDG 16 has the largest coverage gap but a below-median semantic gap, while SDG 3 has a large coverage gap and the second-largest semantic gap). Research attention coexists with both low- and high-gap cases (e.g. SDG 9 high-coverage low-gap, SDG 3 high-coverage high-gap), which is why H1c is near zero. No SDG drives a positive association; the null result is consistent across all configurations.



**Figure 6.** Coverage gap vs semantic gap, one point per SDG.

## 4.5 Cross-Sensitivity Robustness of the Gap Rankings

The semantic gap rankings discussed below are not equally stable across all measurement choices. Table 3 shows each SDG’s gap rank (1 = largest gap) under the canonical configuration (LR, full policy corpus, cap 50) and ten alternatives spanning four axes: embedding architecture (MiniLM 384-d: LR, MLP, zero-shot), assignment method (zero-shot nearest-centroid, MLP), policy source family (curated AI/SDG, SDGi, UNGDC), and segment cap (20, none).

Three tiers emerge. First, SDG 17 is the most method-sensitive SDG: it sits at rank 17 (smallest gap) under LR and rank 16 under MLP but rises to rank 7 under zero-shot assignment. SDG 13 and SDG 11 are the most robustly high-gap goals, staying in the top four under all three assignment methods (LR ranks 1, 3; zero-shot ranks 4, 2; MLP ranks 3, 4). SDG 3 also stays high across methods (LR 2, zero-shot 5, MLP 6). Second, several SDGs shift dramatically by assignment method alone: SDG 12 jumps from rank 7 under LR to rank 1 under MLP; SDG 5 moves from rank 9 under LR to rank 16 under zero-shot; SDG 8 moves from rank 6 to rank 15. Third, across policy-source and cap configurations, SDG 11 stays in the top five under every sub-corpus and cap; SDGs 13 and 3 also stay consistently high (ranges 1–7 and 1–10 respectively). SDG 4 is the most stable low-gap goal, staying in ranks 14–15 across all nine configurations. The discussion limits its claims to findings that survive this filtering.

The same cross-sensitivity structure holds for the coverage gap (Table 4). The within-SDG coverage-gap ranking is stable across embedding architectures ( $\rho = 0.89$  for MiniLM–LR versus the MPNet–LR baseline) and across retrieval strategy ( $\rho = 0.96$  for concept-based field-of-study retrieval versus the keyword baseline), indicating that the coverage conclusions are not driven by the choice of query terms.

## 4.6 Robustness of the Semantic-Gap Ranking to the Distance Functional

The centroid gap (§4.3) is the cosine distance between per-SDG mean embeddings. A reviewer may object that a single mean discards the shape of the two point clouds: two SDGs could share a mean direction yet differ in spread or modality (Villani, 2008). This section tests that objection directly by recomputing the per-SDG research–policy gap under a battery of distribution-aware metrics on the identical point clouds, holding the supervised SDG assignment and embeddings fixed; only the distance functional changes.

The canonical gap is reproduced exactly as a validation gate: the full-corpus centroid gap equals the published value to  $10^{-4}$ , and the streamed research means reproduce the committed centroids exactly. Full-corpus metrics (sliced Wasserstein (Bonneel et al., 2015),

**Table 3.** Cross-sensitivity robustness of within-SDG semantic gap rankings across embedding architectures and assignment methods.

| SDG                  | Encoder (embedding architecture) |           |           |               |           |           | Policy source |           |           | Segment cap |           | Retrieval |
|----------------------|----------------------------------|-----------|-----------|---------------|-----------|-----------|---------------|-----------|-----------|-------------|-----------|-----------|
|                      | mpnet (768d)                     |           |           | MiniLM (384d) |           |           | Curated       | SDGi      | UNGDC     | 20          | None      | Concept   |
|                      | LR                               | MLP       | ZS        | LR            | MLP       | ZS        |               |           |           |             |           |           |
| SDG 1                | <i>4</i>                         | <i>1</i>  | <i>8</i>  | <i>8</i>      | <i>3</i>  | <i>16</i> | <i>3</i>      | <i>4</i>  | <i>7</i>  | <i>4</i>    | <i>4</i>  | <i>1</i>  |
| SDG 2                | <b>13</b>                        | <b>14</b> | <b>11</b> | <b>14</b>     | <b>15</b> | <b>12</b> | <b>11</b>     | <b>13</b> | <b>11</b> | <b>13</b>   | <b>12</b> | <b>11</b> |
| SDG 3                | <b>2</b>                         | <b>5</b>  | <b>5</b>  | <b>2</b>      | <b>2</b>  | <b>8</b>  | <b>9</b>      | <b>2</b>  | <b>1</b>  | <b>1</b>    | <b>6</b>  | <b>4</b>  |
| SDG 4                | <b>15</b>                        | <b>17</b> | <b>14</b> | <b>16</b>     | <b>17</b> | <b>13</b> | <b>16</b>     | <b>14</b> | <b>16</b> | <b>15</b>   | <b>15</b> | <b>14</b> |
| SDG 5                | <i>10</i>                        | <i>15</i> | <i>16</i> | <i>4</i>      | <i>12</i> | <i>14</i> | <i>12</i>     | <i>11</i> | <i>10</i> | <i>10</i>   | <i>9</i>  | <i>9</i>  |
| SDG 6                | <b>14</b>                        | <b>11</b> | <b>7</b>  | <b>15</b>     | <b>13</b> | <b>6</b>  | <b>10</b>     | <b>15</b> | <b>13</b> | <b>14</b>   | <b>14</b> | <b>15</b> |
| SDG 7                | <b>5</b>                         | <b>4</b>  | <b>1</b>  | <b>6</b>      | <b>8</b>  | <b>4</b>  | <b>1</b>      | <b>8</b>  | <b>2</b>  | <b>5</b>    | <b>3</b>  | <b>8</b>  |
| SDG 8                | <b>6</b>                         | <b>7</b>  | <b>15</b> | <b>5</b>      | <b>4</b>  | <b>15</b> | <b>6</b>      | <b>5</b>  | <b>3</b>  | <b>7</b>    | <b>5</b>  | <b>3</b>  |
| SDG 9                | 12                               | 9         | 17        | 10            | 9         | 17        | 14            | 1         | 12        | 11          | 13        | 13        |
| SDG 10               | <i>11</i>                        | <i>10</i> | <i>13</i> | <i>3</i>      | <i>1</i>  | <i>7</i>  | <i>8</i>      | <i>12</i> | <i>14</i> | <i>12</i>   | <i>11</i> | <i>10</i> |
| SDG 11               | <i>3</i>                         | <i>2</i>  | <i>2</i>  | <i>7</i>      | <i>6</i>  | <i>9</i>  | <i>4</i>      | <i>3</i>  | <i>4</i>  | <i>3</i>    | <i>2</i>  | <i>2</i>  |
| SDG 12               | <i>7</i>                         | <i>3</i>  | <i>9</i>  | <i>11</i>     | <i>7</i>  | <i>11</i> | <i>2</i>      | <i>7</i>  | <i>6</i>  | <i>6</i>    | <i>7</i>  | <i>12</i> |
| SDG 13               | <b>1</b>                         | <b>8</b>  | <b>4</b>  | <b>1</b>      | <b>5</b>  | <b>2</b>  | <b>5</b>      | <b>6</b>  | <b>5</b>  | <b>2</b>    | <b>1</b>  | <b>5</b>  |
| SDG 14               | <i>8</i>                         | <i>6</i>  | <i>3</i>  | <i>12</i>     | <i>14</i> | <i>3</i>  | <i>7</i>      | <i>9</i>  | <i>8</i>  | <i>8</i>    | <i>8</i>  | <i>6</i>  |
| SDG 15               | <i>16</i>                        | <i>16</i> | <i>12</i> | <i>9</i>      | <i>16</i> | <i>5</i>  | <i>13</i>     | <i>16</i> | <i>15</i> | <i>16</i>   | <i>16</i> | <i>17</i> |
| SDG 16               | <i>9</i>                         | <i>12</i> | <i>10</i> | <i>13</i>     | <i>11</i> | <i>10</i> | <i>15</i>     | <i>10</i> | <i>9</i>  | <i>9</i>    | <i>10</i> | <i>7</i>  |
| SDG 17               | <b>17</b>                        | <b>13</b> | <b>6</b>  | <b>17</b>     | <b>10</b> | <b>1</b>  | <b>17</b>     | <b>17</b> | <b>17</b> | <b>17</b>   | <b>17</b> | <b>16</b> |
| Rank Corr ( $\rho$ ) | 1.00                             | 0.80      | 0.49      | 0.71          | 0.69      | -0.03     | 0.76          | 0.80      | 0.92      | 0.99        | 0.97      | 0.89      |

*Notes:* Each cell reports the within-SDG semantic gap rank (1 = largest gap, 17 = smallest gap) under that encoder and assignment method. The base encoder is **all-mpnet-base-v2** (768-d); the alternative encoder is **all-MiniLM-L6-v2** (384-d). LR = canonical supervised logistic-regression classifier; policy segments are capped at 50 per source document per SDG (Assumption A-CHUNKCAT). Zero-shot = nearest-centroid assignment on the SDG reference centroids. MLP = 4-layer/384-hidden network retrained on the full training pool. Policy-source and segment-cap columns are LR-based. **Bold** = encoding-invariant (rank difference  $\leq 1$  between MPNet and MiniLM for the same assignment method); *italic* = encoder-sensitive (rank difference  $\geq 4$ ). Rank Corr ( $\rho$ ) is the Spearman correlation of each column’s SDG gap ranks against the canonical MPNet-LR column.

Chamfer (Borgefors, 1988), Gaussian-2-Wasserstein (Dowson & Landau, 1982), Grassmann (Edelman et al., 1998), linear MMD (Gretton et al., 2012)) are computed over all assigned papers and carry no sampling variance; sampled metrics (exact EMD (Rubner et al., 2000), Sinkhorn (Cuturi, 2013), energy (Székely & Rizzo, 2013), RBF-MMD (Gretton et al., 2012), C2ST (Lopez-Paz & Oquab, 2017)) use a 50,000-paper draw per SDG, with SDGs 1 and 17 using their full population.

**Shape-sensitive metrics preserve the ranking.** Gaussian-2-Wasserstein ( $\rho = 0.838$ ), RBF maximum mean discrepancy ( $\rho = 0.627$ ), exact optimal-transport EMD ( $\rho = 0.892$ ), sliced Wasserstein ( $\rho = 0.755$ ), and energy distance ( $\rho = 0.748$ ; (Sejdinovic et al., 2013)) all preserve the canonical ordering at  $p < 0.001$  (Spearman, 1904). SDG 13 is the most divergent SDG under every metric; its primacy is unconditional.

**What centroids discard is quantified, not hidden.** The Gaussian-2-Wasserstein distance decomposes into a mean term and a shape (covariance) term (Dowson & Landau, 1982; Heusel et al., 2017). Across SDGs the shape term accounts for a mean share of  $\approx 56\%$  of the distance. Centroids therefore drop real distributional signal—but that signal

**Table 4.** Cross-sensitivity robustness of within-SDG coverage-gap rankings across embedding architectures and assignment methods.

| SDG                  | Encoder (embedding architecture) |           |           |               |           |           | Policy source |           |           | Retrieval |
|----------------------|----------------------------------|-----------|-----------|---------------|-----------|-----------|---------------|-----------|-----------|-----------|
|                      | mpnet (768d)                     |           |           | MiniLM (384d) |           |           | Curated       | SDGi      | UNGDC     | Concept   |
|                      | LR                               | MLP       | ZS        | LR            | MLP       | ZS        |               |           |           |           |
| SDG 1                | 10                               | 11        | 14        | 15            | 13        | 12        | 17            | 8         | 13        | 10        |
| SDG 2                | 11                               | 14        | 13        | 14            | 15        | 11        | 10            | 11        | 12        | 14        |
| SDG 3                | <b>2</b>                         | <b>2</b>  | <b>2</b>  | <b>2</b>      | <b>2</b>  | <b>3</b>  | <b>3</b>      | <b>2</b>  | <b>3</b>  | <b>2</b>  |
| SDG 4                | <b>6</b>                         | <b>6</b>  | <b>4</b>  | <b>6</b>      | <b>4</b>  | <b>4</b>  | <b>4</b>      | <b>15</b> | <b>6</b>  | <b>7</b>  |
| SDG 5                | <b>15</b>                        | <b>15</b> | <b>15</b> | <b>16</b>     | <b>16</b> | <b>15</b> | <b>11</b>     | <b>5</b>  | <b>16</b> | <b>12</b> |
| SDG 6                | 16                               | 16        | 12        | 13            | 14        | 9         | 16            | 9         | 15        | 17        |
| SDG 7                | <b>8</b>                         | <b>8</b>  | <b>7</b>  | <b>7</b>      | <b>7</b>  | <b>7</b>  | <b>7</b>      | <b>13</b> | <b>8</b>  | <b>9</b>  |
| SDG 8                | <b>13</b>                        | <b>10</b> | <b>17</b> | <b>12</b>     | <b>12</b> | <b>17</b> | <b>12</b>     | <b>4</b>  | <b>9</b>  | <b>11</b> |
| SDG 9                | <b>4</b>                         | <b>5</b>  | <b>5</b>  | <b>5</b>      | <b>5</b>  | <b>6</b>  | <b>1</b>      | <b>1</b>  | <b>4</b>  | <b>4</b>  |
| SDG 10               | 14                               | 17        | 16        | 10            | 11        | 16        | 15            | 17        | 14        | 15        |
| SDG 11               | 7                                | 7         | 9         | 9             | 8         | 14        | 5             | 3         | 7         | 6         |
| SDG 12               | 9                                | 9         | 11        | 11            | 10        | 13        | 13            | 16        | 10        | 8         |
| SDG 13               | <b>5</b>                         | <b>3</b>  | <b>3</b>  | <b>4</b>      | <b>3</b>  | <b>2</b>  | <b>6</b>      | <b>12</b> | <b>5</b>  | <b>5</b>  |
| SDG 14               | <b>17</b>                        | <b>13</b> | <b>10</b> | <b>17</b>     | <b>17</b> | <b>10</b> | <b>14</b>     | <b>14</b> | <b>17</b> | <b>16</b> |
| SDG 15               | 12                               | 12        | 8         | 8             | 9         | 8         | 9             | 10        | 11        | 13        |
| SDG 16               | <b>1</b>                         | <b>1</b>  | <b>6</b>  | <b>1</b>      | <b>1</b>  | <b>5</b>  | <b>8</b>      | <b>7</b>  | <b>1</b>  | <b>1</b>  |
| SDG 17               | <b>3</b>                         | <b>4</b>  | <b>1</b>  | <b>3</b>      | <b>6</b>  | <b>1</b>  | <b>2</b>      | <b>6</b>  | <b>2</b>  | <b>3</b>  |
| Rank Corr ( $\rho$ ) | 1.00                             | 0.94      | 0.79      | 0.89          | 0.91      | 0.70      | 0.79          | 0.35      | 0.96      | 0.96      |

*Notes:* Each cell reports the within-SDG coverage-gap rank (1 = largest gap, 17 = smallest gap) under that encoder and assignment method. The base encoder is **all-mpnet-base-v2** (768-d); the alternative encoder is **all-MiniLM-L6-v2** (384-d). LR = canonical supervised logistic-regression classifier; MLP = 4-layer/384-hidden network; Zero-shot = nearest-centroid assignment on the SDG reference centroids.

Coverage gap = |research proportion – policy proportion| per SDG, using document-weighted policy proportions (Assumption A19) for all methods. Policy-source columns compare the canonical keyword-retrieved research profile against each policy-source family. The Retrieval column replaces keyword retrieval with concept-based (AI/ML field-of-study) retrieval. The Segment-cap axis is omitted: coverage gap is segment-cap-independent. Rank Corr ( $\rho$ ) is the Spearman correlation of each column’s SDG coverage-gap ranks against the canonical MPNet-LR column.

does not reorder the SDGs, which is why the centroid ranking is robust.

**Margin-level nuance.** The second and third canonical positions (SDGs 3 and 11) are less stable: several distribution-aware metrics elevate SDGs 1, 7, 12, and 14 instead. This is a genuine margin-level reshuffle among SDGs with similar gap magnitudes, not an error; the overall monotonic agreement (positive  $\rho$  for every shape-sensitive metric) holds.

**Separation, not ranking, from C2ST.** The classifier two-sample test (Lopez-Paz & Oquab, 2017) returns AUC  $\approx 1$  for all 17 SDGs, so it confirms that research and policy are separable in every SDG but carries no ranking signal ( $\rho \approx 0$ ). Its near-constant value is expected, not a failure.

**Sampling stability.** The sampled metrics use a 50,000-paper draw; a replicate seed changes any sampled gap by at most 0.013, confirming the estimate is stable and the default seed is representative. Together with the cross-sensitivity checks (§4.5), the centroid gap

**Table 5.** Per-SDG research-policy gap under distribution-aware metrics versus the canonical centroid gap (Part 1 of 2: full-corpus exact metrics). SWD = sliced Wasserstein;  $\text{MMD}^2$  = squared RBF maximum mean discrepancy; Gauss.  $W_2$  = Gaussian 2-Wasserstein; C2ST = classifier two-sample test AUC. The final row reports Spearman  $\rho$  between each metric’s ranking and the canonical ranking across the 17 SDGs.

| SDG                  | Canonical | SWD   | Chamfer | Hausdorff | Gauss. $W_2$ | Grassmann | linear MMD |
|----------------------|-----------|-------|---------|-----------|--------------|-----------|------------|
| 1                    | 0.435     | 0.017 | 0.469   | 0.639     | 0.863        | 3.921     | 0.322      |
| 2                    | 0.285     | 0.016 | 0.373   | 0.492     | 0.776        | 3.733     | 0.257      |
| 3                    | 0.464     | 0.017 | 0.411   | 0.554     | 0.811        | 3.637     | 0.295      |
| 4                    | 0.262     | 0.014 | 0.386   | 0.493     | 0.690        | 3.333     | 0.206      |
| 5                    | 0.333     | 0.016 | 0.409   | 0.545     | 0.799        | 3.728     | 0.276      |
| 6                    | 0.282     | 0.016 | 0.404   | 0.520     | 0.767        | 3.714     | 0.258      |
| 7                    | 0.434     | 0.018 | 0.419   | 0.558     | 0.831        | 3.753     | 0.325      |
| 8                    | 0.427     | 0.016 | 0.437   | 0.566     | 0.802        | 3.616     | 0.276      |
| 9                    | 0.289     | 0.014 | 0.389   | 0.530     | 0.738        | 3.475     | 0.200      |
| 10                   | 0.309     | 0.014 | 0.443   | 0.569     | 0.751        | 3.430     | 0.221      |
| 11                   | 0.460     | 0.016 | 0.422   | 0.546     | 0.809        | 3.663     | 0.307      |
| 12                   | 0.425     | 0.018 | 0.471   | 0.645     | 0.856        | 3.674     | 0.326      |
| 13                   | 0.478     | 0.019 | 0.388   | 0.521     | 0.833        | 3.718     | 0.373      |
| 14                   | 0.392     | 0.018 | 0.414   | 0.552     | 0.831        | 3.786     | 0.323      |
| 15                   | 0.229     | 0.014 | 0.407   | 0.522     | 0.702        | 3.471     | 0.195      |
| 16                   | 0.333     | 0.014 | 0.396   | 0.518     | 0.732        | 3.428     | 0.214      |
| 17                   | 0.219     | 0.013 | 0.378   | 0.466     | 0.681        | 3.318     | 0.198      |
| $\rho$ vs. canonical | –         | 0.755 | 0.520   | 0.588     | 0.838        | 0.498     | 0.826      |

ranking is therefore robust to both the assignment method and the distance functional.

**Table 6.** Per-SDG research–policy gap under distribution-aware metrics versus the canonical centroid gap (Part 2 of 2: sampled metrics; continuation of Table 5). EMD = exact 1-Wasserstein; MMD<sup>2</sup> = squared RBF-MMD; Centroid (samp.) = sampled centroid-gap control column.

| SDG                  | Canonical | EMD   | Sinkhorn | Energy | MMD <sup>2</sup> | C2ST AUC | Centroid (samp.) |
|----------------------|-----------|-------|----------|--------|------------------|----------|------------------|
| 1                    | 0.435     | 0.696 | 0.552    | 0.277  | 0.120            | 0.996    | 0.435            |
| 2                    | 0.285     | 0.599 | 0.519    | 0.235  | 0.110            | 0.997    | 0.284            |
| 3                    | 0.464     | 0.683 | 0.562    | 0.250  | 0.110            | 0.996    | 0.464            |
| 4                    | 0.262     | 0.575 | 0.507    | 0.185  | 0.088            | 0.997    | 0.262            |
| 5                    | 0.333     | 0.647 | 0.537    | 0.245  | 0.111            | 0.996    | 0.333            |
| 6                    | 0.282     | 0.590 | 0.531    | 0.240  | 0.115            | 0.998    | 0.282            |
| 7                    | 0.434     | 0.662 | 0.551    | 0.287  | 0.130            | 0.998    | 0.434            |
| 8                    | 0.427     | 0.672 | 0.536    | 0.232  | 0.103            | 0.996    | 0.428            |
| 9                    | 0.289     | 0.621 | 0.523    | 0.176  | 0.079            | 0.997    | 0.288            |
| 10                   | 0.309     | 0.639 | 0.598    | 0.190  | 0.085            | 0.994    | 0.309            |
| 11                   | 0.460     | 0.656 | 0.530    | 0.261  | 0.118            | 0.997    | 0.463            |
| 12                   | 0.425     | 0.695 | 0.593    | 0.283  | 0.124            | 0.998    | 0.424            |
| 13                   | 0.478     | 0.656 | 0.494    | 0.329  | 0.149            | 0.998    | 0.478            |
| 14                   | 0.392     | 0.648 | 0.559    | 0.291  | 0.134            | 0.997    | 0.392            |
| 15                   | 0.229     | 0.576 | 0.532    | 0.178  | 0.084            | 0.998    | 0.230            |
| 16                   | 0.333     | 0.625 | 0.488    | 0.185  | 0.083            | 0.997    | 0.335            |
| 17                   | 0.219     | 0.546 | 0.445    | 0.181  | 0.086            | 0.995    | 0.219            |
| $\rho$ vs. canonical | –         | 0.892 | 0.348    | 0.748  | 0.627            | -0.047   | 1.000            |

## 5 Discussion

This section interprets the findings reported above. It moves from what the data show to what they mean: the structure behind the two divergence dimensions, the robust patterns, and the implications and limits of the study. No new results are introduced here.

The headline result of this study is a structural dissociation: the coverage gap — how differently research and policy prioritise each SDG — and the within-SDG semantic gap — how differently they frame the same SDG — are independent dimensions of divergence, not linked effects. The primary test (H1a) finds no detectable association in the canonical specification ( $r = -0.083$ ,  $p = 0.751$ ;  $n = 17$ ). Given the small number of SDG-level observations, this should be interpreted as evidence against a strong monotonic relationship rather than evidence of exact independence. All four hypotheses (H1a–H1d) return null results (Table 2); the null holds across all sensitivity checks excluding SDG 4 or SDG 17. Coverage and framing should therefore be treated as separate axes of divergence, and one cannot be read as a proxy for the other.

This dissociation is possible because the study separates two methodological roles that are often conflated: classification and measurement. The supervised LR classifier (§3.4) is a labeling instrument that assigns texts to SDGs using a transparent 768-coefficient decision boundary. The frozen Sentence-BERT embedding space (§3.3) is a separate measurement surface that provides a consistent geometry for cross-corpus comparison.



Classification determines which texts belong to each SDG; semantic distance is then computed independently using centroids in the embedding space. Because the two tasks use different procedures, the dissociation between coverage and framing is a finding about the discourse structure of the corpora, not a design confound.

The paper makes two contributions, and they carry different weight. The first is a measurement framework that separates classification from measurement, coverage from framing. This framework holds up across changes of assignment method, policy source family, and segment cap, as the checks in Section 4.5 and the appendices show. The second is a set of findings about how AI research and SDG policy actually relate. These findings are weaker: the relationship between coverage and semantic gaps reverses sign under some configurations (Table 2), and SDG 17’s position depends heavily on the assignment method (Table 3). The framework should be read as the paper’s main contribution. The findings about alignment should be read as a first application, worth checking with a larger and more diverse policy corpus, not as settled results.

The coverage and gap profiles are themselves conditional on corpus composition. The policy source-family split shifts the full-corpus profile: the curated AI/SDG sub-corpus is innovation-heavy (SDG 9 leads), while SDGi and UNGDC are climate-, partnerships-, and governance-oriented (SDGs 13, 17, 16 lead) (Table 13, Appendix H.2). The cross-sensitivity table (Table 3) is therefore the primary interpretive device: it reports gap ranks under all four sensitivity axes, and the discussion limits its claims to findings that survive this filtering.

## 5.1 Two Distinct Dimensions of Observed Divergence

This dissociation is best interpreted as a structural difference between knowledge systems rather than as a simple failure of communication. The Mode 1/2 distinction (Gibbons et al., 1994) sharpens this reading. The research corpus concentrates on technically tractable applications in health, energy, and infrastructure — SDGs 3, 7, and 9 — where dominant ML tasks align with well-defined problem-solving rather than institutional coordination, while the policy-side dominance of SDGs 13, 16, and 17 reflects Mode 2 goal-setting arising from contexts of application that demand coordination and collective action. Epistemic-community theory limits the claim: if expert influence partly depends on shared problem framing under uncertainty (Haas, 1992), then large within-SDG semantic gaps mark where such shared framing is less plausible. Following Cross (2013), however, semantic proximity is only a possible precondition for knowledge transfer; it cannot show whether expert communities exist, whether they have policy access, or whether policy actors use the relevant research. The two-communities theory (Caplan, 1979) reinforces this structural reading: the uncorrelatedness of research coverage and semantic gap is what one would expect if the two corpora represent distinct knowledge cultures with different

reward systems and knowledge-use criteria, rather than two sides of a shared enterprise.

The practical implication is that observed research–policy divergence is a multidimensional problem: coverage and framing must be measured as separate axes, and neither can serve as a proxy for the other.

## 5.2 Robust Patterns of Divergence

The cross-sensitivity comparison narrows the set of gap findings stable enough for individual interpretation. SDG 15 is the robustly lowest-gap SDG under the canonical LR and MLP classifiers (rank 16 LR, rank 17 MLP) but sits at rank 11 under zero-shot. SDG 17’s position shifts dramatically by assignment method: it is rank 17 (smallest gap) under LR, rank 16 under MLP, but rises to rank 7 under zero-shot nearest-centroid — the most method-sensitive SDG in the study. This confirms that partnership discourse was the most fragile measurement point under the centroid-based approach and benefits from the sharper decision boundaries of supervised classifiers. The sensitivity is consistent with SDG 17’s distinctive character as the dedicated means-of-implementation goal, whose policy discourse centres on coordination and resourcing rather than a shared substantive topic (Le Blanc, 2015); a centroid-based method struggles to anchor this qualitatively different semantic boundary, while the supervised LR isolates a tighter region where research and policy texts converge.

SDG 13, SDG 11, and SDG 3 are the most robustly high-gap SDGs, staying in the top six ranks across all three assignment methods (LR, MLP, zero-shot) and across cap and policy-source configurations. These are substantive domains (climate action, sustainable cities, health) where research and policy discourse diverge most sharply. The climate gap (SDG 13) is the most consistent high-gap finding — it ranks first under LR, third under MLP, and fourth under zero-shot, while staying in the top two under all policy sub-corpora and segment caps.

SDG 15 is the robustly closest SDG across all three assignment methods, occupying the bottom two ranks under LR and MLP (rank 16, rank 17) and rank 11 under zero-shot. It does not prove substantive alignment, but it shows the instrument can separate a close case from divergent ones. SDG 9, previously the closest gap under the centroid-based approach, now sits in the middle-to-lower band with a gap of 0.289 — research and policy still frame infrastructure in relatively close terms, but the gap is no longer the narrowest.

The main coverage patterns are themselves robust. The full policy corpus is dominated by the governance–climate–partnerships block (SDGs 16, 13, and 17). Appendix H.2 notes that this is source-family sensitive — the curated AI/SDG sub-corpus is more innovation-oriented — but the broader institutional profile holds across the UNGDC and SDGi families. The research corpus, conversely, concentrates on technically tractable

SDGs (3, 9, 7, 16, 11).

SDGs 7 and 12 are consistently in the upper half of gap ranks across most configurations, while SDG 11 is the most consistently high-gap goal outside the top three (LR 3, zero-shot 2, MLP 4). SDGs 1, 2, 5, 6, 8, 9, 10, 14, and 16 move by four or more ranks between some configurations (e.g. SDG 8 shifts from rank 6 under LR to rank 15 under zero-shot; SDG 9 shifts from rank 13 under LR and MLP to rank 1 under the SDGi-only reference). No individual interpretive claims are made about their precise gap positions.

To make the gap measure concrete, two examples illustrate what high and low semantic distance look like at the term level<sup>1</sup>. SDG 13 (Climate Action), one of the widest gaps (rank 1, gap 0.478), epitomises the high-gap pattern: research texts emphasise *model*, *prediction*, *neural*, and *analysis* — technical measurement language — while policy texts emphasise *climate*, *change*, *emissions*, *countries*, and *action* — institutional commitment language. Both sides concern the same goal, but with fundamentally different orientations: one measures, the other commits. SDG 17 (Partnerships), by contrast, shows the narrowest gap under the supervised classifier (rank 17, gap 0.219): research texts are distinguished by technical terms (*model*, *method*, *machine*) and policy texts by institutional terms (*countries*, *nations*, *global*, *agenda*), yet both sides converge on a shared core of partnership and cooperation vocabulary that keeps the embedding centroids close despite peripheral lexical differences. Together, these cases show that the semantic gap captures a real difference in how the two discourses frame shared SDG assignments, not a uniform property across all goals.

### 5.3 Implications

Two implications follow from the combined findings, and each argues against reading topical mention as alignment.

**Research and policy agenda-setting.** The map does not imply that funding should be redirected on the basis of embedding distances; it shows where research–policy translation and brokerage are most needed. In Guston’s terms, high-gap SDGs may be candidates for boundary-organising practices that create usable boundary objects while remaining accountable to both knowledge producers and policy users (Guston, 2001).

**SDG-aware research communication.** Increasing research output on neglected SDGs is insufficient; research communication norms also matter. Future work could test whether publication incentives that encourage explicit SDG framing improve semantic proximity more effectively than topic-level funding mandates alone. In practice, this could mean

---

<sup>1</sup>Distinctive terms are TF-IDF-weighted unigrams from deterministic samples of research and policy texts assigned to each SDG by the LR classifier. Full details and additional cases are reported in Appendix B.1.

funder-required policy-facing summaries, journal-level SDG justification fields, or policy synthesis briefs. Following Cash et al. (2003), this is boundary work rather than communication alone, requiring accountability to both scientific and policy principals (Guston, 2001).

These implications speak to distinct audiences. For SDG monitoring bodies, the finding that shared labels do not imply shared framing challenges keyword-based SDG mapping: coverage and framing should be reported as separate dimensions. For research funders, the results suggest that funding allocation alone is unlikely to close the semantic gap without also attending to how research communicates its policy relevance; portfolio-level semantic gap tracking could complement keyword-based counts by measuring whether funded research moves toward or away from policy discourse over successive cycles. For journals adopting SDG classification tags, the finding that framing divergence varies by SDG warns against treating SDG tagging as sufficient evidence of policy-aligned research.

## 5.4 Limitations

The cross-sensitivity comparison (Table 3) bounds which gap findings are stable enough to interpret. The limitations below address scope conditions that this comparison does not resolve.

### **Scope: semantic proximity, not substantive alignment or policy influence.**

The study measures semantic proximity in embedding space, not whether research and policy are substantively aligned or whether research is used by policy actors. Embedding proximity reflects shared vocabulary, not shared agenda, and does not measure policy influence, institutional mechanisms, or knowledge use — these would require methods beyond this study’s scope.

**Distance metric choice.** This study uses cosine distance between mean embeddings. A distribution-aware robustness check (Section 4.6) recomputes the per-SDG research–policy gap under Gaussian-2-Wasserstein, exact optimal-transport Earth Mover’s Distance, sliced Wasserstein, energy distance, RBF maximum mean discrepancy, and Chamfer distance on the identical point clouds, holding the assignment fixed. The SDG ranking is preserved (Gaussian-2-Wasserstein  $\rho = 0.838$ , exact EMD  $\rho = 0.892$ , sliced Wasserstein  $\rho = 0.755$ , energy  $\rho = 0.748$ ), and SDG 13 is the most divergent SDG under every metric. Centroids therefore discard real distributional shape — the mean Gaussian-2-Wasserstein distance decomposes into a mean term and a shape (covariance) term, and the shape term accounts for a mean share of  $\approx 56\%$  — but that shape does not reorder the SDGs. The one metric that does not track the ranking, the classifier two-sample test (C2ST), returns  $\text{AUC} \approx 1$  in all SDGs and is a separation test rather than a ranker. Sampling stability is verified with a replicate seed: the maximum change in any sampled gap is 0.013.

**Corpus scope and retrieval boundaries.** The policy corpus represents international institutional SDG discourse rather than the universe of AI sustainability policy. Domestic plans, sub-national policies, and non-English documents are not represented; the research corpus is also English-only. Both biases likely understate Global South perspectives. On the research side, the four-term query prioritises precision over recall, a deliberate choice (Section 3.1) that minimises non-AI noise but may miss specialised AI subfields whose practitioners do not use the four umbrella terms in their titles or abstracts. The OECD.AI methodology (OECD, 2024) addresses this through field-of-study membership (the Artificial Intelligence or Machine Learning field of study in MAG, inherited by OpenAlex) rather than keyword search. The corpus is therefore a high-precision subset rather than a comprehensive census of AI-for-sustainability research. Results are conditional on these boundaries. Appendix H.2 makes the source-family sensitivity explicit.

**OpenAlex upstream classification.** The research corpus is drawn from OpenAlex, whose SDG-topic tagging is itself a downstream classifier trained on external SDG-labelled data; the SDG assignments entering this study are therefore one step from human judgement and inherit that upstream model’s bias. Records that OpenAlex does not tag to any SDG are absent from the retrieval universe, so the corpus reflects OpenAlex’s classification boundary rather than the full population of AI-for-sustainability research. This is a method-dependent corpus boundary in the sense of Armitage et al. (2020), and coverage proportions should be read as conditional on it.

**Corpus asymmetry: AI-for-SDG scope.** As noted in Section 3.1, the research corpus is  $\text{AI} \cap \text{SDG}$  while the policy corpus is closer to a union of institutional SDG discourse and AI governance documents. Between-SDG rankings are conditional on this asymmetric construction.

**Assignment-method sensitivity of gap estimates.** The semantic gap is conditional on how texts are assigned to SDGs. The cross-sensitivity table (Table 3) compares gap rankings under LR argmax, zero-shot nearest-centroid, and the MLP robustness check. The three methods produce varied correlations between their gap values (LR vs. MLP Spearman  $\rho = 0.88$ , LR vs. zero-shot  $\rho = 0.48$ , MLP vs. zero-shot  $\rho = 0.50$ ), and SDG 17 in particular flips from the smallest gap under both supervised classifiers (LR rank 17, MLP rank 16) to a mid-ranked gap under zero-shot (rank 7). The LR-based estimates are the canonical result; the zero-shot and MLP columns are the uncertainty bounds. The raw LR gaps are retained because they best match what the study measures and are most stable across the other two sensitivity axes (policy source family and segment cap). The encoder axis (MPNet 768-d vs MiniLM 384-d) confirms this is not an artefact of the MPNet geometry: within each assignment method the two encoders’ gap rankings agree strongly (LR  $\rho \approx 0.71$ , zero-shot  $\rho \approx 0.69$ ), so the zero-shot column’s near-zero  $\rho$  against the LR baseline ( $\approx -0.03$ ) is a method-axis effect (zero-shot vs. LR), not an encoder one –

MiniLM zero-shot still tracks MPNet zero-shot.

**Register effects on the semantic gap.** Semantic distance may capture differences in communicative function (see Biber and Conrad (2009) on the register–genre–style distinction). Research abstracts and policy documents are text varieties designed to do different things, so the gap is a measure of observed textual-semantic divergence, not a purified estimate of substantive disagreement. Appendix F reports a one-direction register-adjusted sensitivity check: subtracting the strongest linear research–policy axis compresses the gaps modestly (mean change -0.021) but preserves the relative ordering. The spread between raw and adjusted estimates is an uncertainty band around the register-sensitive component. Aggressive multi-direction subtraction risks overcorrecting, and no single adjustment cleanly separates register from substantive topic. The raw gap is retained because it best matches what the study measures.

**Cross-sectional design.** The study observes one period, 2018–2025, against a policy corpus assembled over a similar timeframe; it cannot address whether alignment is improving or deteriorating, or whether specific policy events (e.g., the EU AI Act 2024, IPCC AR6 2022) have shifted research framing.

## 6 Conclusion

This paper introduces a measurement framework that separates two dimensions of research-policy divergence — coverage and semantic framing — and demonstrates empirically that these dimensions are independent, not linked. The dissociation between coverage gaps and within-SDG semantic gaps holds across three assignment methods, four policy-source families, and two segment-cap configurations: topical attention is not a proxy for framing proximity.

The framework’s contribution is portable. It can be applied to any research-policy interface where both sides share a category system and texts can be embedded in a shared semantic space — conservation science, health governance, or technology assessment, for example. The cross-sensitivity structure (Table 3) provides a reusable template for bounding uncertainty in such comparisons: rather than reporting a single gap ranking, the framework makes the stability of each position transparent across the measurement choices that analysts control. The findings about AI-for-sustainability alignment specifically are a first application of this approach, not settled results.

This measurement framework does not solve the alignment problem. It operationalises one dimension of it — semantic proximity in embedding space — and makes the boundary conditions of that operationalisation explicit. Semantic distance cannot tell us whether policy actors use research, whether epistemic communities exist, or whether observed

divergence is productive or problematic. What it can do is show, at scale and with verifiable assumptions, where closer investigation is most warranted.

Alignment should therefore be understood as a multi-level empirical construct — from lexical overlap through semantic proximity to substantive agreement (Section 2.2) — rather than as a binary property inferred from shared labels. The framework developed here advances the measurement from Level 1 to Level 2. Level 3 remains the domain of qualitative, institutional, and deliberative methods that this study does not attempt.

## References

- Armitage, C. S., Lorenz, M., & Mikki, S. (2020). Mapping scholarly publications related to the sustainable development goals: Do independent bibliometric approaches get the same results? *Quantitative Science Studies*, 1(3), 1092–1108. [https://doi.org/10.1162/qss\\_a\\_00071](https://doi.org/10.1162/qss_a_00071)
- Bautista-Puig, N., Aleixo, A. M., Leal, S., Azeiteiro, U., & Costas, R. (2021). Unveiling the research landscape of sustainable development goals and their inclusion in higher education institutions and research centers: Major trends in 2000–2017. *Frontiers in Sustainability*, 2, 620743. <https://doi.org/10.3389/frsus.2021.620743>
- Bexell, M., & Jönsson, K. (2017). Responsibility and the united nations' sustainable development goals. *Forum for Development Studies*, 44(1), 13–29. <https://doi.org/10.1080/08039410.2016.1252424>
- Biber, D. (1988). *Variation across speech and writing*. Cambridge University Press.
- Biber, D., & Conrad, S. (2009). *Register, genre, and style*. Cambridge University Press.
- Bonneel, N., Rabin, J., Peyré, G., & Papadakis, N. (2015). Sliced and radon wasserstein barycenters of measures. *Journal of Mathematical Imaging and Vision*, 51, 22–45. <https://doi.org/10.1007/s10851-014-0506-3>
- Borgefors, G. (1988). Hierarchical chamfer matching: A parametric edge matching algorithm. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 10(6), 849–865. <https://doi.org/10.1109/34.9108>
- Bornmann, L. (2017). Measuring impact in research evaluations: A thorough discussion of methods for, effects of and problems with impact measurements. *Higher Education*, 73, 775–787. <https://doi.org/10.1007/s10734-016-9995-x>
- Bornmann, L., Haunschild, R., & Marx, W. (2016). Policy documents as sources for measuring societal impact: How often is climate change research mentioned in policy-related documents? *Scientometrics*, 109(3), 1477–1495. <https://doi.org/10.1007/s11192-016-2115-y>
- Caplan, N. (1979). The two-communities theory and knowledge utilization. *American Behavioral Scientist*, 22(3), 459–470. <https://doi.org/10.1177/000276427902200308>

- Cash, D. W., Clark, W. C., Alcock, F., Dickson, N. M., Eckley, N., Guston, D. H., Jäger, J., & Mitchell, R. B. (2003). Knowledge systems for sustainable development. *Proceedings of the National Academy of Sciences*, 100(14), 8086–8091. <https://doi.org/10.1073/pnas.1231332100>
- Cowls, J., Tsamados, A., Taddeo, M., & Floridi, L. (2021). A definition, benchmark and database of AI for social good initiatives. *Nature Machine Intelligence*, 3(2), 111–115. <https://doi.org/10.1038/s42256-021-00296-0>
- Cross, M. K. D. (2013). Rethinking epistemic communities twenty years later. *Review of International Studies*, 39(1), 137–160. <https://doi.org/10.1017/S0260210512000034>
- Cuturi, M. (2013). Sinkhorn distances: Lightspeed computation of optimal transport [arXiv:1306.0895]. In C. J. Burges, L. Bottou, M. Welling, Z. Ghahramani, & K. Weinberger (Eds.), *Advances in neural information processing systems 26 (nips 2013)* (pp. 2292–2300). Curran Associates, Inc.
- Dowson, D. C., & Landau, B. V. (1982). The fréchet distance between multivariate normal distributions. *Journal of Multivariate Analysis*, 12(3), 450–455. [https://doi.org/10.1016/0047-259X\(82\)90077-X](https://doi.org/10.1016/0047-259X(82)90077-X)
- Edelman, A., Arias, T. A., & Smith, S. T. (1998). The geometry of algorithms with orthogonality constraints. *SIAM Journal on Matrix Analysis and Applications*, 20(2), 303–353. <https://doi.org/10.1137/S0895479895290954>
- Fukuda-Parr, S. (2016). From the millennium development goals to the sustainable development goals: Shifts in purpose, concept, and politics of global goal setting for development. *Gender & Development*, 24(1), 43–52. <https://doi.org/10.1080/13552074.2016.1145895>
- Gibbons, M., Limoges, C., Nowotny, H., Schwartzman, S., Scott, P., & Trow, M. (1994). *The new production of knowledge: The dynamics of science and research in contemporary societies*. SAGE Publications.
- Gjorgjevikj, A., Mishev, K., Trajanov, D., & Kocarev, L. (2025). Benchmarking sentence encoders in associating indicators with sustainable development goals and targets. *IEEE Access*, 13, 141434–141460. <https://doi.org/10.1109/ACCESS.2025.3595894>
- Gretton, A., Borgwardt, K. M., Rasch, M. J., Schölkopf, B., & Smola, A. J. (2012). A kernel two-sample test. *Journal of Machine Learning Research*, 13, 723–773.
- Griggs, D. J., Nilsson, M., Stevance, A., & McCollum, D. (Eds.). (2017). *A guide to SDG interactions: From science to implementation*. International Council for Science. <https://doi.org/10.24948/2017.01>
- Guston, D. H. (2001). Boundary organizations in environmental policy and science: An introduction. *Science, Technology, & Human Values*, 26(4), 399–408. <https://doi.org/10.1177/016224390102600401>



- Haas, P. M. (1992). Introduction: Epistemic communities and international policy coordination. *International Organization*, 46(1), 1–35. <https://doi.org/10.1017/S0020818300001442>
- Harris, Z. S. (1954). Distributional structure. *WORD*, 10(2–3), 146–162. <https://doi.org/10.1080/00437956.1954.11659520>
- Harvard Dataverse. (2024). United nations general debate corpus [Sessions 1–80 (1946–2025)]. <https://dataverse.harvard.edu/dataset.xhtml?persistentId=doi:10.7910/DVN/0TJX8Y>
- Heilinger, J.-C., Kempt, H., & Nagel, S. K. (2024). Beware of sustainable AI! uses and abuses of a worthy goal. *AI and Ethics*, 4(2), 201–212. <https://doi.org/10.1007/s43681-023-00259-8>
- Heink, U., Marquard, E., Heubach, K., Jax, K., Kugel, C., Neßhöver, C., Neumann, R. K., Paulsch, A., Tilch, S., Timaeus, J., & Vandewalle, M. (2015). Conceptualizing credibility, relevance and legitimacy for evaluating the effectiveness of science-policy interfaces: Challenges and opportunities. *Science and Public Policy*, 42(5), 676–689. <https://doi.org/10.1093/scipol/scu082>
- Hellwig, J., Huggett, S., & Siebert, M. (2019). *Artificial intelligence: How knowledge is created, transferred, and used*. Elsevier.
- Heras-Saizarbitoria, I., Urbieto, L., & Boiral, O. (2022). Organizations’ engagement with sustainable development goals: From cherry-picking to SDG-washing? *Corporate Social Responsibility and Environmental Management*, 29(2), 316–328. <https://doi.org/10.1002/csr.2202>
- Heusel, M., Ramsauer, H., Unterthiner, T., Nessler, B., & Hochreiter, S. (2017). Gans trained by a two time-scale update rule converge to a local nash equilibrium [arXiv:1706.08500]. *Advances in Neural Information Processing Systems 30 (NIPS 2017)*.
- Hickel, J. (2019). The contradiction of the sustainable development goals: Growth versus ecology on a finite planet. *Sustainable Development*, 27(5), 873–884. <https://doi.org/10.1002/sd.1947>
- Kashnitsky, Y., Roberge, G., Mu, J., Kang, K., Wang, W., Vanderfeesten, M., Rivest, M., Chamezopoulos, S., Jaworek, R., Vignes, M., Jayabalasingham, B., Boonen, F., James, C., Doornenbal, M., & Labrosse, I. (2024). Evaluating approaches to identifying research supporting the united nations sustainable development goals. *Quantitative Science Studies*, 5(2), 408–425. [https://doi.org/10.1162/qss\\_a\\_00304](https://doi.org/10.1162/qss_a_00304)
- Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2012). Imagenet classification with deep convolutional neural networks. In F. Pereira, C. Burges, L. Bottou, & K. Weinberger (Eds.), *Advances in neural information processing systems* (Vol. 25). Curran Associates, Inc. [https://proceedings.neurips.cc/paper\\_files/paper/2012/file/c399862d3b9d6b76c8436e924a68c45b-Paper.pdf](https://proceedings.neurips.cc/paper_files/paper/2012/file/c399862d3b9d6b76c8436e924a68c45b-Paper.pdf)

- Le Blanc, D. (2015). Towards integration at last? the sustainable development goals as a network of targets. *Sustainable Development*, 23(3), 176–187. <https://doi.org/10.1002/sd.1582>
- Lopez-Paz, D., & Oquab, M. (2017). Revisiting classifier two-sample tests [arXiv:1610.06545]. *Proceedings of the 5th International Conference on Learning Representations (ICLR)*. <https://doi.org/10.48550/arXiv.1610.06545>
- McCarthy, C., Brooks, C., Sternberg, T., Shaney, K., & Hoshino, B. (2026). Ai-assisted multi-target classification for research-policy alignment in conservation science. *Ecological Informatics*, 94, 103669. <https://doi.org/10.1016/j.ecoinf.2026.103669>
- Nedungadi, P., Surendran, S., Tang, K.-Y., & Raman, R. (2024). Big data and AI algorithms for sustainable development goals: A topic modeling analysis. *IEEE Access*, 12, 188519–188541. <https://doi.org/10.1109/ACCESS.2024.3516500>
- Nilsson, M., Griggs, D., & Visbeck, M. (2016). Policy: Map the interactions between sustainable development goals. *Nature*, 534(7607), 320–322. <https://doi.org/10.1038/534320a>
- OECD. (2024). Oecd.ai: Partner data and methodological note [OECD.AI Observatory methodological documentation]. <https://oecd.ai/en/partner-data-methodological-note>
- Priem, J., Piwowar, H., & Orr, R. (2022). OpenAlex: A fully-open index of the world’s scholarly works. <https://doi.org/10.48550/arXiv.2205.01833>
- Pukelis, L., Bautista-Puig, N., Statulevičiūtė, G., Stančiauskas, V., Dikmener, G., & Akylbekova, D. (2022). OSDG 2.0: A multilingual tool for classifying text data by UN sustainable development goals (SDGs). *arXiv preprint*. <https://doi.org/10.48550/arXiv.2211.11252>
- Reimers, N., & Gurevych, I. (2019). Sentence-BERT: Sentence embeddings using siamese BERT-networks. *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*, 3982–3992. <https://doi.org/10.18653/v1/D19-1410>
- Rodriguez, P. L., Spiraling, A., & Stewart, B. M. (2023). Embedding regression: Models for context-specific description and inference. *American Political Science Review*, 117(4), 1255–1274. <https://doi.org/10.1017/S0003055422001228>
- Rubner, Y., Tomasi, C., & Guibas, L. J. (2000). The earth mover’s distance as a metric for image retrieval. *International Journal of Computer Vision*, 40(2), 99–121. <https://doi.org/10.1023/A:1026543900054>
- Schoeberl, M., Toney, B., & Dunham, M. (2023). *Identifying ai research* (tech. rep.). Center for Security and Emerging Technology (CSET). <https://doi.org/10.51593/20230>
- SDG Classification Expert Group. (2025). SDG classification benchmark [Expert-labelled benchmark for SDG text classification, all 17 SDGs]. <https://github.com/SDGClassification/benchmark>

- SDGi. (2024). SDGi corpus of voluntary national and local reviews [Available on Hugging Face]. <https://huggingface.co/datasets/UNDP/sdgi-corpus>
- Sejdinovic, D., Sriperumbudur, B., Gretton, A., & Fukumizu, K. (2013). Equivalence of distance-based and rkhs-based statistics in hypothesis testing. *The Annals of Statistics*, 41(5), 2263–2291. <https://doi.org/10.1214/13-AOS1140>
- Sentence-Transformers. (2021). all-mpnet-base-v2: Sentence embedding model.
- Singh, A., Kanaujia, A., Singh, V. K., & Vinuesa, R. (2024). Artificial intelligence for sustainable development goals: Bibliometric patterns and concept evolution trajectories. *Sustainable Development*, 32(1), 724–754. <https://doi.org/10.1002/sd.2706>
- Sioumalas-Christodoulou, K., & Tympas, A. (2025). AI metrics and policymaking: Assumptions and challenges in the shaping of AI. *AI & Society*, 40, 4655–4670. <https://doi.org/10.1007/s00146-025-02181-5>
- Spearman, C. (1904). The proof and measurement of association between two things. *The American Journal of Psychology*, 15(1), 72–101. <https://doi.org/10.2307/1412159>
- Strauss, I., Moure, I., O'Reilly, T., & Rosenblat, S. (2025, April). Real-world gaps in AI governance research: AI safety and reliability in everyday deployments. <https://doi.org/10.48550/arXiv.2505.00174>
- Székely, G. J., & Rizzo, M. L. (2013). Energy statistics: A class of statistics based on distances. *Journal of Statistical Planning and Inference*, 143(8), 1249–1272. <https://doi.org/10.1016/j.jspi.2013.03.018>
- Toney-Wails, A., Curlee, K., & Probasco, E. (2024). Trust issues: Discrepancies in trustworthy AI keywords use in policy and research. *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency (FAccT '24)*, 2222–2233. <https://doi.org/10.1145/3630106.3659035>
- Vanderfeesten, M., Spielberg, E., & Gunes, Y. (2020). Survey data of “Mapping Research Output to the Sustainable Development Goals (SDGs)” by Aurora universities network (AUR) [Dataset, version 1.0.1]. <https://doi.org/10.5281/zenodo.3798385>
- Villani, C. (2008). *Optimal transport: Old and new* (Vol. 338). Springer. <https://doi.org/10.1007/978-3-540-71050-9>
- Vinuesa, R., Azizpour, H., Leite, I., Balaam, M., Dignum, V., Domisch, S., Felländer, A., Langhans, S. D., Tegmark, M., & Fuso Nerini, F. (2020). The role of artificial intelligence in achieving the sustainable development goals. *Nature Communications*, 11(1), 233. <https://doi.org/10.1038/s41467-019-14108-y>
- Wang, K., Shen, Z., Huang, C., Wu, C.-H., Dong, Y., & Kanakia, A. (2020). Microsoft academic graph: When experts are not enough. *Quantitative Science Studies*, 1(1), 396–413. [https://doi.org/10.1162/qss\\_a\\_00021](https://doi.org/10.1162/qss_a_00021)
- Wulff, D. U., Meier, D. S., & Mata, R. (2023). SDG Knowledge Hub dataset of SDG-labeled news articles [IISD SDG Knowledge Hub; CC-BY 4.0; 9,172 articles]. <https://doi.org/10.5281/zenodo.7523032>

## A Diagnostics for Reference Classifier

### A.1 SDG 4 Lexical Artefact Audit

The SDG 4 concern is not used to relabel research records. Its narrower purpose is to test whether the concern is empirically plausible. Table 7 therefore compares the lexical composition of SDG 4-assigned research records with a matched non-SDG4 sample and with SDG 9, using only two transparent keyword groups: ML-learning vocabulary and education-specific vocabulary.

The audit points to a more nuanced conclusion than a pure false-positive story. More than half of the SDG 4-assigned set contains education-specific vocabulary without ML-only terms, and roughly another third contains both education and ML-learning vocabulary. By contrast, the matched non-SDG4 sample and the SDG 9 comparison set are dominated by ML-only vocabulary. This means the SDG 4 signal is not simply the same technical language that drives SDG 9. However, the large “both” share also confirms that ML-learning language and education language overlap materially within the SDG 4-assigned set.

The main analysis therefore retains the canonical LR argmax assignment but continues to treat SDG 4 as artefact-affected. The most accurate reading is a learning-related assignment whose interpretation is blurred by lexical overlap — neither “clean education” nor “pure ML artefact.”


**Table 7.** SDG 4 lexical artefact audit.

| Corpus subset            | N      | ML-only % | Education-only % | Both % | Neither / unclear % |
|--------------------------|--------|-----------|------------------|--------|---------------------|
| SDG 4-assigned research  | 453225 | 10.5      | 45.3             | 23.9   | 20.4                |
| Non-SDG4 research sample | 453225 | 50.3      | 4.1              | 2.7    | 42.9                |
| SDG 9-assigned research  | 566759 | 58.8      | 2.7              | 2.9    | 35.6                |

*Notes:* The audit is not used to reassign SDGs. It only checks whether SDG 4-assigned research is disproportionately dominated by machine-learning vocabulary, education vocabulary, or both.

### A.2 Centroid Similarity

This subsection reports the pairwise cosine-similarity heatmap (Figure 7) for the 17 canonical SDG centroids. The centroids are the per-SDG mean training embeddings (L2-normalised) used by the LR classifier. The heatmap provides supplementary geometric context for the classification geometry that Section 3.4 builds on.



**Figure 7.** Pairwise cosine similarity between canonical SDG centroids (lower triangle). SDG 13 and SDG 17 sit close together (cosine 0.820), as do SDGs 1, 8, and 10.

Figure 7 shows the full 17×17 similarity grid. The lower triangle is annotated; the upper triangle is left blank because the matrix is symmetric. Reading the grid, three observations stand out. First, the diagonal dominates: most centroids are far more similar to themselves than to any neighbour, which is why the supervised LR classifier (Table 1) can build on this well-separated base to reach 0.823. Second, a small set of off-diagonal cells is elevated. SDG 11’s centroid neighbours SDG 9 (cosine 0.745), and SDG 8 lies near SDGs 1 and 10 (cosine 0.806 and 0.863), and SDGs 1 and 10 are themselves the closest off-diagonal pair (cosine 0.930) — these are the proximities that make the SDG 11 and SDG 8 scores in Section 4.1 indicative rather than exact. Third, SDG 17 is the clear outlier: its row and column are the flattest, and its nearest neighbour is SDG 13 (cosine 0.820), reflecting the overlap between climate and partnership discourse noted in Section 4.1.

## B Supplementary Robustness and Sensitivity Checks

### B.1 Lexical Illustration of the Semantic Gap

Table 8 adds a small lexical interpretability check for two high-gap, policy-dominant cases (SDGs 17 and 13) and one lower-gap comparison case (SDG 9). The table is generated from deterministic samples of existing hard-assigned research and policy texts. It reports distinctive terms rather than selected quotations, so it is a descriptive guide to language patterns, not qualitative coding or proof of policy intent.

The contrast helps make the embedding distances more legible. In SDGs 17 and 13, the high semantic gaps correspond to a split between technical research language and more institutional, international, national-government, and climate-action policy language. SDG 17 is especially cautionary: its research-side terms are biomedical/ML-heavy, while its policy-side terms are international and agenda-oriented. SDG 9, by contrast, shows more overlap around technology, infrastructure, innovation, and systems language. This does not prove substantive alignment in SDG 9 or substantive failure in SDGs 13 and 17; it simply illustrates why shared SDG labels can conceal different textual framings.

### B.2 Implications for Interpreting the Main Estimates

These diagnostics are methodological stress tests that clarify the limits of the LR classifier approach, rather than failed repairs. The PCA diagnostics show that real-world SDG discourse is fuzzy rather than cleanly clustered, especially in the research corpus. The SDG 4 audit shows that the concern is real, but more nuanced than a pure machine-learning false positive. The policy source-family split shows that the full-corpus policy profile is a broader institutional SDG discourse rather than a universal AI policy profile. The canonical hard-threshold classifier assignment ( $\text{argmax}$ ) is retained because it is transparent, reproducible, and directly interpretable.

These analyses also show that the LR classifier should be interpreted as a transparent semantic anchoring procedure rather than a definitive operational SDG system. Attempts to address multi-label overlap and corpus-register bias improve some aspects of the design but introduce new modelling choices. In this paper, the hard-threshold classifier assignment is therefore retained as the canonical specification, while the alternatives are reported as robustness and boundary-condition checks.

**Table 8.** Lexical interpretability check for selected semantic-gap cases.

| SDG    | Gap   | Research-side distinctive terms                                         | Policy-side distinctive terms                                                           | Descriptive reading                                                                                                                                                                      |
|--------|-------|-------------------------------------------------------------------------|-----------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| SDG 17 | 0.219 | learning, analysis, model, dan, yang, method, machine, used             | countries, nations, oecd, united, national, agenda, climate, country                    | High-gap case: the research-side sample is technical and biomedical/ML-heavy, while policy language is international, institutional, and agenda-oriented.                                |
| SDG 13 | 0.478 | model, learning, prediction, method, machine, neural, accuracy, amp     | climate, change, emissions, countries, national, action, global, nations                | High-gap case: both sides concern climate, but policy language is more institutional and commitment-oriented while research language is more technical and measurement-oriented.         |
| SDG 9  | 0.289 | learning, model, neural, method, machine, algorithm, proposed, analysis | infrastructure, national, government, innovation, population, public, countries, sector | Lower-gap comparison: differences remain visible, but the contrast is closer to a technology/infrastructure register split than to the broader institutional gap seen in SDGs 13 and 17. |

*Notes:* Terms are generated from deterministic samples of research and policy texts assigned to each SDG. The table is descriptive only and is not used to relabel records or replace the embedding-based semantic-gap estimates.

## C Sample-Stability Robustness Check

To test whether the headline findings depend heavily on research-corpus size, I reran the downstream research-side analysis on repeated random subsamples drawn from the pre-embedded, pre-scored OpenAlex corpus. For each of 11 sampled tiers (from 1,000 to 2,000,000 papers), I drew 100 independent samples without replacement, recomputed the research coverage profile, rebuilt sampled research centroids, recalculated the mean within-SDG semantic gap, and re-estimated the policy-text calibration-bias summary. Policy-side quantities were held fixed. The full 3,105,144-paper analysis is shown as a deterministic anchor row rather than as another repeated sample.

*Plain-language guide.*<sup>2</sup> The stability ladder shows that the clearest practical plateau

<sup>2</sup>The three metrics are: (1) mean SD of SDG coverage shares — how much the topic balance moves across repeated samples; (2) mean semantic gap — how different research and policy language are within the same SDG; (3) policy–research score gap — how much the scoring instrument favours policy text over research text regardless of content.

**Table 9.** Sample-stability ladder for the research corpus.

| Sample Size | Mean SD of SDG coverage shares | Mean within-SDG semantic gap | Policy-text calibration bias |
|-------------|--------------------------------|------------------------------|------------------------------|
| 1,000       | 0.006                          | $0.427 \pm 0.020$            | $0.112 \pm 0.008$            |
| 2,000       | 0.004                          | $0.394 \pm 0.014$            | $0.113 \pm 0.006$            |
| 5,000       | 0.003                          | $0.371 \pm 0.009$            | $0.112 \pm 0.004$            |
| 10,000      | 0.002                          | $0.365 \pm 0.006$            | $0.113 \pm 0.003$            |
| 20,000      | 0.001                          | $0.361 \pm 0.005$            | $0.113 \pm 0.002$            |
| 50,000      | 0.001                          | $0.358 \pm 0.003$            | $0.113 \pm 0.001$            |
| 100,000     | 0.001                          | $0.357 \pm 0.002$            | $0.113 \pm 0.001$            |
| 200,000     | 0.000                          | $0.357 \pm 0.002$            | $0.113 \pm 0.001$            |
| 500,000     | 0.000                          | $0.357 \pm 0.001$            | $0.113 \pm 0.000$            |
| 1,000,000   | 0.000                          | $0.356 \pm 0.001$            | $0.113 \pm 0.000$            |
| 2,000,000   | 0.000                          | $0.356 \pm 0.000$            | $0.113 \pm 0.000$            |
| Full corpus | —                              | 0.356                        | 0.113                        |

*Notes:* Mean SD of SDG coverage shares is the unweighted mean of the per-SDG standard deviations in research coverage share across the 100 random draws at each sampled tier. The remaining columns report the mean and draw-to-draw standard deviation of the headline downstream metrics. The full-corpus row is deterministic and therefore has no variance term.

appears by roughly 200,000 papers, where the mean semantic gap ( $\sim 0.357$ ) and policy-text calibration bias ( $\sim 0.113$ ) already stabilise to within 0.001 of their full-corpus values. From 500,000 papers upward, draw-to-draw dispersion is negligible. The full 3,105,144-paper result is retained as the primary estimate because it uses all available data while delivering the tightest values, but the substantive conclusions are not artefacts of sample size choice. The larger corpus also enables per-SDG subgroup detail and tighter rank-order precision that the plateau minimum does not guarantee.

## D Model Selection: Grid Search Protocol and Rationale

### D.1 Models Tested

Three model families were evaluated under 5-fold GroupKFold cross-validation (document-level grouping, not random splits) on the pooled training data ( $n=52,779$ ). The test set ( $n=9,734$ ) was never touched during any grid search.

- **Logistic Regression (LR):** Interpretable per-dimension coefficients per SDG; fast; standard multiclass baseline.
- **MLP (PyTorch):** Can capture non-linear interactions in the embedding space.



- **RF / XGB:** Considered but excluded because the embedding space is already a learned representation — bagged trees add ensemble complexity without a clear mechanism for improvement over LR/MLP, and would not produce interpretable coefficients per dimension.

## D.2 LR Grid Search

The LR grid searched  $C \in \{100, 10, 1, 0.1, 0.01\}$  with L2 penalty. The best configuration was  $C=10.0$ ,  $\text{penalty}=l2$ ,  $\text{class\_weight}=\text{None}$ ,  $\text{solver}=lbfgs$  ( $\text{macro-F1} = 0.8107 \pm 0.0066$  CV). An extended grid adding L1 penalty ( $l1\_ratio=1$ ) and balanced  $\text{class\_weight}$  found no improvement: L1 was  $6\times$  slower and scored 0.8066, while balanced  $\text{class\_weight}$  scored 0.8070. The embedding space is dense — L1 sparsity offers no benefit — and the dual correction of weighted loss plus macro-averaged scoring degrades majority-class performance without compensating minority classes.  $C=10.0$  and  $C=1.0$  are essentially tied (difference  $\sim 0.001$ ), but  $C=10.0$  was selected for conservatism.

## D.3 MLP Grid Search

The MLP swept  $n\_layers \in \{1, 2, 4, 8, 16\} \times \text{hidden\_size} \in \{256, 384\}$  (20 configurations), followed by a learning-rate sweep ( $lr \in \{1e-4, 3e-4, 1e-3, 3e-3\}$ ) on the best architectures. The best MLP (4 layers, 384 hidden,  $lr=3e-4$ ) achieved  $\text{macro-F1} = 0.8243 \pm 0.0058$ . Learning rate had negligible impact — all combinations within  $\sim 0.003$  of the best.

## D.4 Unified Ranking and Final Selection

| Rank  | Config                | CV macro-F1         |
|-------|-----------------------|---------------------|
| 1     | MLP 4L/384h           | $0.8243 \pm 0.0058$ |
| 2     | MLP 4L/256h           | $0.8242 \pm 0.0066$ |
| 3–15  | MLP (2L–4L variants)  | 0.819–0.824         |
| 16–26 | MLP (8L–16L variants) | 0.817–0.821         |
| 27    | LR $C=10$ , L2        | $0.8107 \pm 0.0066$ |
| 28    | LR $C=1$ , L2         | $0.8093 \pm 0.0065$ |

**Champion: LR ( $C=10.0$ , L2).** The gap between LR and MLP is 0.0137 macro-F1 ( $\sim 1.5$  pooled  $\sigma$ ), a modest difference insufficient to justify the complexity trade-off. MLP requires 4 hidden layers, AdamW, learning rate tuning, early stopping, and dropout selection (36 configurations swept). LR requires one free parameter ( $C$ ), no hidden layers, no dropout, no random seed sensitivity. For this study’s purpose — building an interpretable supervised reference to compare corpora, not a classifier built to label single documents — LR is the right choice. LR provides 768 signed coefficients per SDG class, each directly attributable

to a specific semantic dimension; MLP’s hidden layers diffuse this information across non-linear transformations. The macro-F1 numbers are reported honestly: LR  $0.811 \pm 0.007$ , MLP  $0.824 \pm 0.006$ , gap  $\sim 1.5\sigma$ . LR is chosen for transparency, not superior accuracy.

## D.5 Held-Out Test Evaluation of the MLP Robustness Check

The MLP test macro-F1 (0.832) cited in Section 4.1 is a separate evaluation from the CV grid-search results above. After the CV sweep, the champion MLP architecture (4 layers, 384 hidden, lr=3e-4) was retrained on the full training pool (n=52,779) and evaluated once against the held-out test set (n=9,734), the same sanctioned single-use test evaluation applied to the LR champion. The resulting test macro-F1 of 0.832 is not directly comparable to the CV mean of  $0.8243 \pm 0.0058$  from the grid search: the CV number is an average across five document-grouped folds on the training pool alone, while the test number is a single evaluation on a disjoint held-out sample after retraining on the full pool. Both numbers confirm that the LR—MLP gap is modest and that architecture choice is not the binding constraint on classification quality.

## E Embedding-Model Sensitivity

### E.1 Scope

An embedding-model sensitivity check compares the canonical **all-mpnet-base-v2** results against the lighter **all-MiniLM-L6-v2** encoder (384-d). This is now performed: the cross-sensitivity table (Table 3) carries both encoders as a fourth axis, and the supervised LR and MLP rankings prove stable across architectures (within-method Spearman  $\rho \approx 0.7$ ), confirming the gap findings are not an artefact of the MPNet geometry. Re-embedding the full research corpus ( $\sim 3.1$ M papers) with a second model required the additional storage noted above, and the sample-stability and cross-sensitivity checks (Sections 4.5 and C) bound robustness under the chosen encoder. The remaining caveat is that MPNet (768-d) and MiniLM (384-d) differ in both architecture and dimensionality, so the comparison bounds encoder sensitivity only modulo the dimensionality drop; a same-dimension encoder pair would be needed to isolate architecture alone. The zero-shot column’s near-zero  $\rho$  against the LR baseline ( $\approx -0.03$ ) is a method-axis effect (zero-shot vs. LR), not an encoder artefact – MiniLM zero-shot still agrees with MPNet zero-shot ( $\rho \approx 0.69$ ). The results are conditional on the MPNet architecture and should be read with that scope in mind.

## F Register-Adjustment Sensitivity

This appendix reports an exploratory register-adjustment procedure that probes whether broad research-policy register differences contribute to the raw semantic gap. Following Biber and Conrad’s distinction between register, genre, and style (Biber & Conrad, 2009), the research-policy contrast is treated primarily as a register contrast: the two corpora differ in situation of use, audience, institutional setting, and communicative purpose, not merely in surface wording. A learned research-policy direction in embedding space may therefore capture a mixture of modality, abstraction, institutional stance, and substantive theme. The adjusted results are reported as a sensitivity diagnostic, not as corrected estimates.

**Procedure.** One global research-versus-policy direction is estimated in the shared embedding space using a binary logistic-regression classifier trained on a balanced sample of 79,987 embeddings (approximately 40,000 per class). On the training set, this classifier achieves accuracy 0.981, showing that a strong corpus-level separation axis exists. This degree of separability is expected for two large corpora drawn from distinct institutional discourse communities and does not entail that the within-SDG gaps are primarily register-driven. Let  $\mathbf{g}$  denote the learned coefficient direction (L2-normalised). For each embedding  $\mathbf{x}$ , the orthogonal projection onto  $\mathbf{g}$  is subtracted:

$$\mathbf{x}^\perp = \mathbf{x} - \frac{\mathbf{x} \cdot \mathbf{g}}{\|\mathbf{g}\|^2} \mathbf{g},$$

and the residual is renormalised to unit length. The result is described as *register-adjusted*, not “debiased”: the subtraction removes one estimated global corpus axis, not every possible register or corpus-composition difference.

For the research side, raw per-SDG means are recovered as  $\mathbf{c}_j^{\text{res}} = \text{research\_centroid}_j \times \text{research\_cohesion}_j$  (since cohesion equals the norm of the raw mean for L2-normalised embeddings), adjusted, and renormalised. For the policy side, the per-SDG means are computed directly from individual adjusted policy embeddings. Adjusted gaps are then  $1 - \cos(\text{adj\_research\_centroid}_j, \text{adj\_policy\_centroid}_j)$ .

**Results.** The high separability between research and policy texts is expected: they represent distinct institutional genres with different audiences, communicative purposes, and stylistic conventions. The logistic-regression classifier’s training accuracy (0.981) captures this genre-level distinction, not a measurement of whether within-SDG gaps are real. Importantly, register adjustment reduces rather than eliminates semantic gaps, and the ranking of SDGs remains stable after adjustment. This suggests that register contributes to observed divergence but does not account for it fully. On average, the

mean gap compresses from 0.356 to 0.335 (a change of -0.021), but the relative ordering of high-gap and low-gap SDGs is preserved: SDGs 13, 11, 1, 7, and 3 remain the furthest apart, while SDGs 15, 4, and 17 remain the closest. The adjustment therefore bounds the register-sensitive component of the gap without replacing the raw estimate.

This single-direction subtraction is a limited diagnostic. It attenuates one broad corpus-level axis but does not cleanly separate register from substantive divergence. The residual gap after adjustment is not a “pure” substantive gap; it only shows what the gap looks like after removing the strongest single linear discriminant between the two corpora. The results are best interpreted as an uncertainty band around the register-sensitive component of the measurement.

**Table 10.** Global one-direction register sensitivity check for the semantic gap. Raw gap = the main semantic-gap estimate from the paper. Adjusted gap = the same within-SDG gap after subtracting one learned global research-policy direction from each embedding.

| SDG    | Description                             | Raw gap | Adjusted gap | $\Delta$ |
|--------|-----------------------------------------|---------|--------------|----------|
| SDG 13 | Climate Action                          | 0.4777  | 0.4558       | -0.0218  |
| SDG 3  | Good Health and Well-Being              | 0.4644  | 0.4509       | -0.0136  |
| SDG 11 | Sustainable Cities and Communities      | 0.4604  | 0.4470       | -0.0134  |
| SDG 1  | No Poverty                              | 0.4354  | 0.4345       | -0.0009  |
| SDG 7  | Affordable and Clean Energy             | 0.4341  | 0.4064       | -0.0277  |
| SDG 8  | Decent Work and Economic Growth         | 0.4270  | 0.4134       | -0.0136  |
| SDG 12 | Responsible Consumption and Production  | 0.4249  | 0.3986       | -0.0263  |
| SDG 14 | Life Below Water                        | 0.3920  | 0.3721       | -0.0198  |
| SDG 16 | Peace, Justice and Strong Institutions  | 0.3335  | 0.3074       | -0.0261  |
| SDG 5  | Gender Equality                         | 0.3330  | 0.3250       | -0.0081  |
| SDG 10 | Reduced Inequalities                    | 0.3091  | 0.2969       | -0.0123  |
| SDG 9  | Industry, Innovation and Infrastructure | 0.2891  | 0.2535       | -0.0356  |
| SDG 2  | Zero Hunger                             | 0.2849  | 0.2568       | -0.0281  |
| SDG 6  | Clean Water and Sanitation              | 0.2817  | 0.2508       | -0.0309  |
| SDG 4  | Quality Education                       | 0.2622  | 0.2232       | -0.0391  |
| SDG 15 | Life on Land                            | 0.2295  | 0.1921       | -0.0373  |
| SDG 17 | Partnerships for the Goals              | 0.2188  | 0.2132       | -0.0057  |
| Mean   |                                         | 0.3563  | 0.3352       | -0.0212  |

## G Concept-Retrieval Sensitivity

Advisor review flagged the four-term keyword query used to retrieve the AI research corpus as an intuitive choice. To test sensitivity to that choice, retrieval was repeated using field-of-study membership for the two AI fields of study (OECD, 2024), mirroring the OECD.AI Observatory methodology, which treats a paper as AI if tagged with the Artificial Intelligence or Machine Learning field of study in the Microsoft Academic Graph (MAG), the taxonomy that OpenAlex later inherited and extended (Priem et al., 2022; Wang et al., 2020). No SDG filter was applied. SDG assignment was performed by the same trained logistic-regression classifier used throughout (Section 3.4), so the only difference from the main corpus is retrieval strategy. The resulting 50,000-paper corpus was passed through the identical segmentation, embedding, scoring, and analysis pipeline.

Table 11 compares the SDG coverage shares of the keyword-retrieved and concept-retrieved corpora. The two profiles are closely aligned (Spearman  $\rho = 0.936$ , Kendall  $\tau = 0.868$ ); the largest deviations are SDG 4 ( $-7.6\text{pp}$ , the known ML-vocabulary artefact, Appendix A.1) and SDG 9 ( $+6.4\text{pp}$ ). The within-SDG coverage-gap and semantic-gap rankings are also stable against the keyword baseline (coverage-gap rank  $\rho = 0.96$ ; semantic-gap rank  $\rho = 0.89$ , Table 3). Schoeberl et al. (2023) independently find that AI-research identification is sensitive to the retrieval method and recommend supervised classification. The present LR scoring is that recommended step applied to SDG assignment. We therefore report the main findings as robust to the query-term choice, while noting that both corpora remain conditional on their respective retrieval strategies.

**Table 11.** SDG coverage-share comparison: keyword-retrieved versus concept-retrieved (AI/ML field-of-study) research corpus.  $\Delta$  is concept minus keyword share in percentage points.

| SDG                                              | Keyword % | Concept % | $\Delta$ |
|--------------------------------------------------|-----------|-----------|----------|
| SDG 1                                            | 0.5       | 0.6       | +0.0     |
| SDG 2                                            | 2.5       | 2.1       | -0.5     |
| SDG 3                                            | 25.3      | 25.9      | +0.6     |
| SDG 4                                            | 14.6      | 7.0       | -7.6     |
| SDG 5                                            | 2.5       | 1.9       | -0.6     |
| SDG 6                                            | 1.1       | 1.0       | -0.1     |
| SDG 7                                            | 4.0       | 4.5       | +0.5     |
| SDG 8                                            | 2.4       | 2.8       | +0.4     |
| SDG 9                                            | 18.3      | 24.6      | +6.4     |
| SDG 10                                           | 1.4       | 0.9       | -0.5     |
| SDG 11                                           | 4.8       | 6.5       | +1.7     |
| SDG 12                                           | 2.1       | 3.8       | +1.7     |
| SDG 13                                           | 4.2       | 5.4       | +1.2     |
| SDG 14                                           | 2.0       | 1.8       | -0.2     |
| SDG 15                                           | 2.6       | 2.6       | +0.0     |
| SDG 16                                           | 11.3      | 8.2       | -3.1     |
| SDG 17                                           | 0.4       | 0.5       | +0.0     |
| Spearman $\rho = 0.936$ ; Kendall $\tau = 0.868$ |           |           |          |

## H Supplementary Cross-Method Data

### H.1 Cross-Method Coverage and Semantic Gap Values

Table 12 reports the raw coverage gap and semantic gap values for each SDG across six encoder–assignment-method combinations, with an additional concept-retrieval LR column for comparison against the keyword-retrieved research corpus baseline (Appendix G). While Tables 3 and 4 rank these values to compare cross-method stability, the table below shows the unrounded gap magnitudes, allowing the rank separation to be interpreted in absolute terms. Coverage gap is document-weighted (Assumption A19); semantic gap policy centroids use a 50-segment-per-document cap.

### H.2 Detailed Policy Profiles by Source Family

The full policy corpus combines a curated AI/SDG sub-corpus, SDGi VNR/VLR reports, and UN General Debate speeches. These are all policy-facing texts, but they are not interchangeable genres. The family split reported in Table 13 therefore asks a narrow question: whether the full-corpus policy profile is broadly stable across source families or whether it is strongly shaped by corpus composition.

The answer is mixed. The broad international-institutional profile is clearly source-family

**Table 12.** Coverage gap and semantic gap values across embedding architectures, assignment methods, and retrieval strategies.

| SDG | Coverage gap (%) |      |      |               |      |      |         | Semantic gap (cosine) |       |       |               |       |       |         |
|-----|------------------|------|------|---------------|------|------|---------|-----------------------|-------|-------|---------------|-------|-------|---------|
|     | MPNet (768d)     |      |      | MiniLM (384d) |      |      | Concept | MPNet (768d)          |       |       | MiniLM (384d) |       |       | Concept |
|     | LR               | MLP  | ZS   | LR            | MLP  | ZS   |         | LR                    | MLP   | ZS    | LR            | MLP   | ZS    |         |
| 1   | 1.5              | 1.6  | 1.6  | 0.8           | 1.2  | 2.7  | 1.5     | 0.435                 | 0.583 | 0.367 | 0.343         | 0.458 | 0.216 | 0.745   |
| 2   | 1.3              | 0.6  | 2.4  | 1.0           | 0.9  | 3.3  | 0.8     | 0.285                 | 0.276 | 0.333 | 0.298         | 0.304 | 0.304 | 0.410   |
| 3   | 23.1             | 21.2 | 22.9 | 24.8          | 24.4 | 20.1 | 23.7    | 0.464                 | 0.451 | 0.425 | 0.482         | 0.468 | 0.403 | 0.596   |
| 4   | 12.4             | 12.8 | 15.3 | 11.5          | 12.2 | 13.7 | 4.9     | 0.262                 | 0.265 | 0.274 | 0.279         | 0.283 | 0.280 | 0.353   |
| 5   | 0.6              | 0.6  | 1.4  | 0.7           | 0.2  | 1.4  | 1.2     | 0.333                 | 0.274 | 0.248 | 0.462         | 0.335 | 0.263 | 0.492   |
| 6   | 0.3              | 0.3  | 2.5  | 1.2           | 1.0  | 3.4  | 0.2     | 0.282                 | 0.283 | 0.416 | 0.291         | 0.314 | 0.429 | 0.345   |
| 7   | 2.8              | 3.3  | 7.0  | 3.7           | 3.4  | 7.2  | 3.2     | 0.434                 | 0.463 | 0.515 | 0.393         | 0.377 | 0.466 | 0.515   |
| 8   | 1.1              | 1.6  | 0.1  | 1.5           | 1.4  | 0.6  | 1.5     | 0.427                 | 0.431 | 0.255 | 0.425         | 0.453 | 0.256 | 0.620   |
| 9   | 15.3             | 13.6 | 10.2 | 12.4          | 12.1 | 8.3  | 21.7    | 0.289                 | 0.315 | 0.219 | 0.331         | 0.368 | 0.210 | 0.383   |
| 10  | 1.0              | 0.2  | 0.5  | 1.8           | 1.5  | 0.8  | 0.5     | 0.309                 | 0.286 | 0.284 | 0.470         | 0.477 | 0.420 | 0.436   |
| 11  | 3.2              | 5.7  | 3.5  | 2.4           | 2.8  | 2.2  | 4.9     | 0.460                 | 0.521 | 0.457 | 0.357         | 0.400 | 0.379 | 0.670   |
| 12  | 1.7              | 2.2  | 2.6  | 1.5           | 1.5  | 2.4  | 3.4     | 0.425                 | 0.469 | 0.355 | 0.318         | 0.398 | 0.347 | 0.402   |
| 13  | 14.6             | 20.0 | 16.8 | 15.9          | 22.1 | 23.9 | 13.4    | 0.478                 | 0.419 | 0.438 | 0.491         | 0.445 | 0.542 | 0.569   |
| 14  | 0.0              | 1.0  | 3.0  | 0.2           | 0.1  | 3.3  | 0.2     | 0.392                 | 0.440 | 0.457 | 0.311         | 0.309 | 0.480 | 0.555   |
| 15  | 1.1              | 1.4  | 3.9  | 3.5           | 2.5  | 6.4  | 1.2     | 0.229                 | 0.272 | 0.314 | 0.335         | 0.300 | 0.430 | 0.313   |
| 16  | 23.7             | 24.7 | 9.5  | 32.4          | 28.3 | 12.9 | 26.8    | 0.333                 | 0.282 | 0.335 | 0.308         | 0.344 | 0.367 | 0.518   |
| 17  | 22.8             | 17.0 | 47.5 | 16.6          | 11.9 | 32.5 | 22.8    | 0.219                 | 0.278 | 0.423 | 0.253         | 0.360 | 0.612 | 0.344   |

*Notes:* Coverage gap = |research proportion – policy proportion| per SDG in percentage points, using document-weighted policy proportions (Assumption A19) for all methods. Semantic gap =  $1 - \cos(\mathbf{r}_{\text{SDG}}, \mathbf{p}_{\text{SDG}})$  where  $\mathbf{r}_{\text{SDG}}$  and  $\mathbf{p}_{\text{SDG}}$  are the research and policy centroids; policy centroids use a 50-segment-per-document cap. All semantic gap values are cosine units.

sensitive. The curated AI/SDG family is much more innovation-oriented, with SDG 9 leading its document-weighted profile, while the SDGi and UNGDC families are more strongly concentrated on climate, partnerships, governance, and, in the SDGi case, SDG 11-style institutional review language. At the same time, the broader full-corpus pattern is not driven by one family alone: both the UNGDC speeches and the full corpus are dominated by the climate–partnerships–governance block, and SDG 17 remains the largest semantic-gap case across all three families.

The implication is interpretive rather than corrective. The full policy corpus should be read as international institutional SDG discourse, not as a universal profile of all AI sustainability policy. The scope of what the policy-side distribution represents is accordingly narrower, but the main diagnostic purpose remains valid.

**Table 13.** Per-SDG policy source-family sensitivity: coverage share and semantic gap.

| SDG    | Full Corpus |              | Curated AI/SDG |              | SDGi VNR/VLR |              | UNGDC      |              |
|--------|-------------|--------------|----------------|--------------|--------------|--------------|------------|--------------|
|        | pol.% (n)   | sem. (n)     | pol.% (n)      | sem. (n)     | pol.% (n)    | sem. (n)     | pol.% (n)  | sem. (n)     |
| 1      | 2.0 (50)    | 0.44 (1,634) | 1.1 (1)        | 0.45 (327)   | 3.8 (12)     | 0.44 (1,188) | 1.8 (37)   | 0.56 (119)   |
| 2      | 1.3 (31)    | 0.29 (1,486) | 0.0 (0)        | 0.32 (239)   | 1.9 (6)      | 0.28 (1,176) | 1.2 (25)   | 0.44 (71)    |
| 3      | 2.2 (54)    | 0.47 (2,428) | 6.4 (6)        | 0.35 (652)   | 12.5 (39)    | 0.52 (1,715) | 0.4 (9)    | 0.68 (61)    |
| 4      | 2.2 (53)    | 0.26 (2,059) | 0.0 (0)        | 0.20 (280)   | 14.4 (45)    | 0.28 (1,712) | 0.4 (8)    | 0.39 (67)    |
| 5      | 3.1 (75)    | 0.33 (1,982) | 0.0 (0)        | 0.31 (271)   | 8.9 (28)     | 0.34 (1,512) | 2.3 (47)   | 0.46 (199)   |
| 6      | 0.8 (20)    | 0.28 (1,214) | 0.0 (0)        | 0.34 (146)   | 2.6 (8)      | 0.28 (1,040) | 0.6 (12)   | 0.41 (28)    |
| 7      | 1.2 (30)    | 0.43 (1,573) | 1.1 (1)        | 0.51 (368)   | 3.5 (11)     | 0.42 (1,138) | 0.9 (18)   | 0.60 (67)    |
| 8      | 1.3 (32)    | 0.43 (1,900) | 0.0 (0)        | 0.42 (315)   | 9.3 (29)     | 0.44 (1,563) | 0.1 (3)    | 0.58 (22)    |
| 9      | 2.9 (72)    | 0.29 (3,004) | 54.3 (51)      | 0.25 (1,871) | 3.5 (11)     | 0.53 (1,107) | 0.5 (10)   | 0.42 (26)    |
| 10     | 0.4 (10)    | 0.31 (1,183) | 0.0 (0)        | 0.37 (311)   | 1.3 (4)      | 0.31 (820)   | 0.3 (6)    | 0.41 (52)    |
| 11     | 1.6 (40)    | 0.46 (1,599) | 0.0 (0)        | 0.44 (155)   | 12.1 (38)    | 0.47 (1,434) | 0.1 (2)    | 0.58 (10)    |
| 12     | 0.4 (9)     | 0.43 (1,125) | 0.0 (0)        | 0.48 (168)   | 2.2 (7)      | 0.43 (953)   | 0.1 (2)    | 0.57 (4)     |
| 13     | 18.8 (462)  | 0.48 (3,077) | 7.4 (7)        | 0.43 (601)   | 4.8 (15)     | 0.44 (1,278) | 21.5 (440) | 0.58 (1,198) |
| 14     | 2.0 (50)    | 0.39 (1,278) | 0.0 (0)        | 0.38 (226)   | 2.2 (7)      | 0.39 (850)   | 2.1 (43)   | 0.52 (202)   |
| 15     | 1.5 (36)    | 0.23 (1,446) | 0.0 (0)        | 0.28 (283)   | 3.5 (11)     | 0.23 (1,090) | 1.2 (25)   | 0.40 (73)    |
| 16     | 35.0 (860)  | 0.33 (4,755) | 8.5 (8)        | 0.21 (967)   | 7.3 (23)     | 0.36 (1,402) | 40.5 (829) | 0.49 (2,386) |
| 17     | 23.3 (571)  | 0.22 (3,792) | 21.3 (20)      | 0.20 (894)   | 6.1 (19)     | 0.20 (1,355) | 26.0 (532) | 0.31 (1,543) |
| $\rho$ | 1.00        | 1.00         | 0.57           | 0.77         | 0.58         | 0.82         | 0.65       | 0.93         |

*Notes:* For each policy source family, **pol.% (n)** is the document-weighted policy share (%) with the document count in parentheses; **sem. (n)** is the within-SDG research-policy semantic gap with the capped segment count in parentheses. Read across a row to see whether the gap is stable or driven by one family.

## I Declaration of AI Use

In accordance with the University’s guidance on generative AI (“Use of GenAI is permitted to support assignments”), this appendix declares the use of AI tools throughout the dissertation. All use was conducted with human oversight and control; AI-generated outputs were reviewed, modified, and independently verified before inclusion. No AI-generated text was submitted without such verification. This declaration follows the order of permitted activities specified in the assessment guidelines. AI-assisted activities not explicitly listed (e.g., code debugging, plotting, LaTeX troubleshooting, literature summarisation) fall under the same oversight and verification standards described in the relevant rows above.



**Table 14.** Declaration of AI use by permitted activity.

| Activity                                    | AI tool(s)                                                                               | How output was used                                                                                                                                                                                                                                                                                                            |
|---------------------------------------------|------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Researching and refining ideas              | ChatGPT, MiniMax / DeepSeek (via OpenCode), GPT (via Codex)                              | ChatGPT was used iteratively to develop ideas that aligned with the author’s intent. Coding agents independently checked data and code availability to confirm feasibility. All final decisions and conclusions are the author’s own. Final research design was formally presented to and approved by supervisor.              |
| Information retrieval / background research | Gemini, MiniMax / DeepSeek (via OpenCode), GPT (via Codex)                               | AI tools read the manuscript for context and recommended papers to investigate further. All cited sources were independently verified against the original publications.                                                                                                                                                       |
| Drafting an outline / organising thoughts   | MiniMax / DeepSeek (via OpenCode), GPT (via Codex)                                       | Generated code scaffolding and assisted with the full analysis pipeline (embedding, scoring, validation, robustness, appendix outputs). The centroid-assignment step was verified against hand-labeled samples, and the scoring and coverage-gap functions were unit-tested. Analytical decisions were directed by the author. |
| Refining research questions                 | ChatGPT, MiniMax / DeepSeek (via OpenCode), GPT (via Codex)                              | Overlaps with row 1. ChatGPT sharpened framing based on reviewed literature; coding agents confirmed questions were answerable with available data and models. The author retained full control over the final formulation. RQs presented to and approved by supervisor.                                                       |
| Checking spelling and grammar               | ChatGPT, Gemini, Claude (web/mobile), MiniMax / DeepSeek (via OpenCode), GPT (via Codex) | Coding agents provided inline stylistic and phrasing suggestions during drafting. Web and mobile versions of ChatGPT, Gemini, and Claude were used for editorial review and proofreading. All suggestions were accepted or rejected per standard editorial judgment.                                                           |